

Database Systems

Office Location: Room 101 opposite lab1 first floor CS building.

Email: anam.gureshi@nu.edu.pk

Mid 1: 15marks

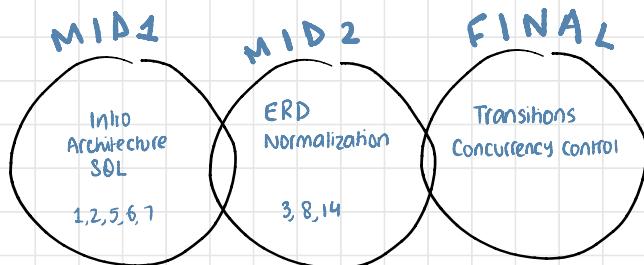
Mid 2: 15 marks

Assignments: 10 marks → 5 assignments

Project: 10 marks

Final Exam: 50 marks

↓
best 4 selected



COURSE PROJECT

- ↳ All Theory + Lab concepts
- ↳ Frontend is connected to backend
- ↳ Frontend design doesn't matter
- ↳ Can pick templates

Database

- ↳ a collection of related data

Entity Relations

- ↳ Faculty ↳ Student take courses
- ↳ Courses ↳ faculty assigned sections
- ↳ Students ↳
- ↳ Sections ↳

Student → Table

S-ID	S-name	S-Ph no.	S-course
------	--------	----------	----------

Mini-World

- ↳ part of real world data stored in a database
student grades at university

Data

- ↳ known facts

Salary
\$ 15 K
\$ 13 K
\$ 13 K
\$ 16 K
random numbers

Database System

- ↳ DBMS software + data itself + applications

entities/relations/objects/table

SAME THING

DBMS (Database management system)

- ↳ software for creating/maintaining database

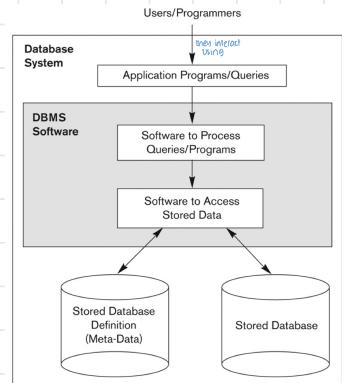
1. Data Manipulation (Query, transaction, modify) ^{insert/delete}
2. Data catalogue
3. Cache → loads whole database to secondary storage medium
4. security and Privacy → prevent unauthorized access
5. Access Control → which data has access
6. Concurrency → each user sees a diff view of db
supports multiple views
7. Presentation and Visualization of data

Applications interact with database by generating

Query

Transaction

- ↳ only read
 - ↳ read and modify/write
 - ↳ fetches data
 - ↳ atomic in nature → execute fully or not at all → COMMIT OR ROLL BACK
- * Applications don't allow unauthorized users to access data



Simple Database

COURSE

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	04	King
92	CS1310	Fall	04	Anderson
102	CS3320	Spring	05	Knuth
112	MATH2410	Fall	05	Chang
119	CS1310	Fall	05	Anderson
135	CS3380	Fall	05	Stone

GRADE_REPORT

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

PREREQUISITE

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310

Shankant B. Navathe

Simplified Database Catalog

RELATIONS

Relation_name	No_of_columns
STUDENT	4
COURSE	4
SECTION	5
GRADE_REPORT	3
PREREQUISITE	2

COLUMNS

Column_name	Data_type	Belongs_to_relation
Name	Character (30)	STUDENT
Student_number	Character (4)	STUDENT
Class	Integer (1)	STUDENT
Major	Major_type	STUDENT
Course_name	Character (10)	COURSE
Course_number	XXXXXXXXXX	COURSE
....
....
....
Prerequisite_number	XXXXXXXXXX	PREREQUISITE

Note: Major_type is defined as an enumerated type with all known majors. XXXXXXXX is used to define a type with four alpha characters followed by four digits

copyright © 2016 Ramez Elmasri and Shamkant B. Navathe

Example of a Database
(with a Conceptual Data Model)

- Mini-world for the example:
 - Part of a UNIVERSITY environment.
- Some mini-world entities:
 - STUDENTS
 - COURSES
 - SECTIONS (of COURSES)
 - (academic) DEPARTMENTS
 - INSTRUCTORS
- Some mini-world relationships:
 - SECTIONS are of specific COURSES
 - STUDENTS take SECTIONS
 - COURSES have prerequisite COURSES
 - INSTRUCTORS teach SECTIONS
 - COURSES are offered by DEPARTMENTS
 - STUDENTS major in DEPARTMENTS

MAIN CHARACTERISTICS OF DATABASE APPROACH

Database Catalog

- ↳ description of a database
 - meta data
- ↳ type of attribute
 - e.g. unique, required
- ↳ constant of attribute
- ↳ data structures
- ↳ constraints

Multiple View Support

- ↳ each user sees a different view of db
- ↳ shows data only which is of interest to that user

Program data Independence

- ↳ insulation b/w programs and data
- ↳ allows changing
 - ↳ data structures
 - ↳ storage organisation
- without changing the DBMS access programs

Data Abstraction

- ↳ a data model is used
- ↳ to hide storage details
- ↳ present users a conceptual view

Sharing of data and multi-user transaction processing:

- Allowing a set of concurrent users to retrieve from and to update the database.
- Concurrency control within the DBMS guarantees that each transaction is correctly executed or aborted
- Recovery subsystem ensures each completed transaction has its effect permanently recorded in the database
- OLTP (Online Transaction Processing) is a major part of database applications. This allows hundreds of concurrent transactions to execute per second.

Database Users

Actors on the scene

- ↳ use and control the db content
- ↳ design, develop and maintain db applications

Workers behind the scene

- those who
 - design and develop DBMS software and computer system operators

1. Database Administrators

- ↳ authorise access to db
- ↳ acquire hardware and software resources
- ↳ monitor efficiency of operations

1. **System Designers and Implementors:** Design and implement DBMS packages in the form of modules and interfaces and test and debug them. The DBMS must interface with applications, language compilers, operating system components, etc.
2. **Tool Developers:** Design and implement software systems called tools for modeling and designing databases, performance monitoring, prototyping, test data generation, user interface creation, simulation etc. that facilitate building of applications and allow using database effectively.
3. **Operators and Maintenance Personnel:** They manage the actual running and maintenance of the database system hardware and software environment.

2. Database Designers

- ↳ define the content, structure, constraints, functions, transactions
- ↳ communicate with end users *→ to understand their needs*

3. End Users

- ↳ use data for queries, reports, update db content

1. Casual

- ↳ access db when needed

2. Naïve / Parametric

- makes up large section of end user*
- ↳ canned transactions *→ previously well defined functions*
- ↳ mostly users of
 - ↳ mobile apps
 - ↳ bank tellers / reservation clerks
 - ↳ social media users *→ post and read info on websites*

5.

▪ Stand-alone:

- Mostly maintain personal databases using ready-to-use packaged applications.
- An example is the user of a tax program that creates its own internal database.
- Another example is a user that maintains a database of personal photos and videos.

4. Sophisticated

- ↳ many use tools in the form of software packages
- ↳ mostly users are
 - ↳ business analysts
 - ↳ scientists, engineers
 - ↳ people familiar with system capabilities

6.

▪ System Analysts and Application Developers

This category currently accounts for a very large proportion of the IT work force.

7.

- System Analysts:** They understand the user requirements of naïve and sophisticated users and design applications including canned transactions to meet those requirements.

8.

- Application Programmers:** Implement the specifications developed by analysts and test and debug them before deployment.

9.

- Business Analysts:** There is an increasing need for such people who can analyze vast amounts of business data and real-time data ("Big Data") for better decision making related to planning, advertising, marketing etc.

Adv OF USING DB Approach

sharing of same data among multiple users

- ↳ Controlling redundancy in data storage ↗
- ↳ restrict unauthorised access to data ↗ privilege commands
- ↳ Provide persistent storage for program objects ↗
- ↳ Provide storage structures ↗ for efficient query processing
- ↳ Provide optimization of queries ↗ for efficient processing
- ↳ Backup and recovery services ↗
- ↳ Multiple interfaces ↗ different view for different class of users
- ↳ Represent complex relationships among data ↗
- ↳ Enforce integrity constraints on db ↗
- ↳ Drawing inferences and actions ↗

When NOT TO USE DBMS

- Main inhibitors (costs) of using a DBMS:
 - High initial investment and possible need for additional hardware.
 - Overhead for providing generality, security, concurrency control, recovery, and integrity functions.
- When a DBMS may be unnecessary:
 - If the database and applications are simple, well defined, and not expected to change.
 - If access to data by multiple users is not required.
- When a DBMS may be infeasible:
 - In embedded systems where a general purpose DBMS may not fit in available storage
- When no DBMS may suffice:

- ↳ If stringent real-time requirements ↗ DBMS overheads a problem ↗ telephone switching systems
- ↳ Complex data ↗ as DBMS has modeling limitations → protion db
- ↳ Special operations ↗ GIS { location based service } ↗ not supported by DBMS

Database Systems

1. Data integrity and consistency ↗ enforces data integrity using constraints and validation
2. Advance search capabilities ↗ filtering
3. Concurrency → multiple users can access same data
4. Better support for large and complex datasets
5. Built in security features ↗ robust security features
6. Redundancy → data normalization
7. Efficient retrieval ↗ SQL queries and indexing
8. Scalability → sharding, replication facilities
9. Program data independence / insulation
10. Self describing nature
11. Multiple views ↗ student and teacher both have different view

Filebase System

1. Simplicity and ease of setup
2. Greater control ↗ directly manage files and folders
3. Better performance for specific tasks ↗ with less processing
4. Lower overhead and resource usage ↗ good for constraints
5. Portability to move/distribute data ↗ when access different systems
6. Simplified backup process ↗ copy files to another location

DATA MODEL

- ↳ what is structure of data
↳ restrictions on valid data
- ↳ what constraints to be applied on data
- ↳ what operations to be applied on data
 - create / read
 - delete
 - search by condition in filter

CONSTRUCTS

- ↳ define db structure
- ↳ includes
 1. elements → and their data type
 2. groups of elements → entity/record/table
 3. relationships

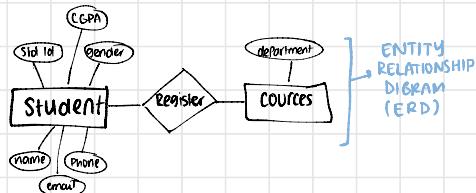
DATA MODEL OPERATIONS

- ↳ used to specify db retrievals and updates
- ↳ includes
 1. basic model operations → insert/delete/update
 2. user defined operations → compute_opn
update_invntry

TYPES OF DATA MODEL

CONCEPTUAL DATA MODEL

- ↳ user friendly → easily understood by everyone



CONVERTED TO

IMPLEMENTATION DATA MODEL

- ↳ conceptual + physical data model

create table Student(

```

    StdId int primary key
    S_name varchar(40)
    S_gender char(1)
    S_phone varchar(13) NOT NULL
    S_add varchar(40)
);
  
```

is automatically NOT NULL, UNIQUE
constraint-type def

PHYSICAL DATA MODEL

- ↳ all about data structures
- ↳ how data is stored
- ↳ DT tree is used to store data in a storage medium

SELF DESCRIBING DATA MODEL

- ↳ at a time key and value both explained
- ↳ providing attributes and data at the same time

```

    Select * Students from University
    where StdId = "1234" AND CGPA > 2.0
  
```

VS

Database Schema

- ↳ description of database
- ↳ includes
- ↳ data types
- ↳ constraints on db
- ↳ description of db structure

aka INTENSION
better to not change once made

Database State/Instance

- ↳ content of a db at a particular moment in time
- ↳ includes collection of all data in db

Schema Diagram

- ↳ illustrative display of db schema

Initial Database State

- ↳ when db is initially loaded in system

Valid State

- ↳ when constraints and structure of db is satisfied

Schema Construct

- ↳ a component of schema
- ↳ STUDENT, COURSE

DISTINCTION

- ↳ db schema changes very infrequently
- ↳ db state changes everytime db is updated

Example of a Database Schema

STUDENT

Name	Student_number	Class	Major
------	----------------	-------	-------

Schema diag
database in

COURSE

Course_name	Course_number	Credit_hours	Department
-------------	---------------	--------------	------------

PREREQUISITE

Course_number	Prerequisite_number
---------------	---------------------

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
--------------------	---------------	----------	------	------------

GRADE_REPORT

Student_number	Section_identifier	Grade
----------------	--------------------	-------

Example of a database state

COURSE

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	04	King
92	CS1310	Fall	04	Anderson
102	CS3320	Spring	05	Knuth
112	MATH2410	Fall	05	Chang
119	CS1310	Fall	05	Anderson
135	CS3380	Fall	05	Stone

GRADE_REPORT

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

PREREQUISITE

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310

Figure 1.2
A database that stores student and course information.

3 schema Architecture

↳ defines DBMS schema at 3 levels

1. Internal Schema ↳ uses physical data model

- ↳ at internal level
- ↳ to describe physical storage structures
- ↳ to access paths

2. Conceptual Schema ↳ uses conceptual/implementation data model

- ↳ at conceptual level
- ↳ describe structure and constraints

3. External Schema ↳ uses conceptual data model

- ↳ at external level
- ↳ describe various user views

↳ SUPPORT 2 characteristics

1. Program data Independence

↳ logical

↳ Physical

↳ change conceptual schema
w/o changing external schema

↳ change internal schema

w/o changing conceptual schema

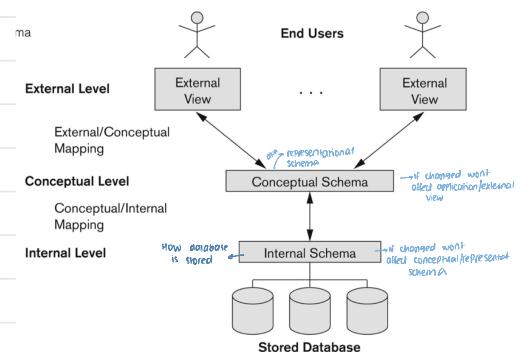
- For example, the internal schema may be changed when certain file structures are reorganized or new indexes are created to improve database performance

2. Multiple views

- ↳ diff views for diff users
- ↳ view data from virtual table

↳ table won't be created
but you can view data
from table

The three-schema architecture



if database schema changed
it will reflect application/external layer

Three-Schema Architecture

- Mappings among schema levels are needed to transform requests and data.
- Programs refer to an external schema, and are mapped by the DBMS to the internal schema for execution.
- Data extracted from the internal DBMS level is reformatted to match the user's external view (e.g. formatting the results of an SQL query for display in a Web page)

Data Independence

- When a schema at a lower level is changed, only the mappings between this schema and higher-level schemas need to be changed in a DBMS that fully supports data independence.
- The higher-level schemas themselves are **unchanged**.
 - Hence, the application programs need not be changed since they refer to the external schemas.

DBMS LANGUAGES

Data Definition language (DDL)

- ↳ used by db designers → to specify conceptual schema

↳ coding which starts with create
table ↗ cursor
view ↗ database

Data manipulation language (DML)

- ↳ used to specify db retrievals and updates

↳ coding operations
create ↗ delete
read ↗ search
in condition
in filter

TYPES OF DML

NON Procedural language

- ↳ specifies what data to retrieve
- ↳ aka high level or declarative

- ↳ no functions
- ↳ no loops
- ↳ no conditions

Procedural language

- ↳ retrieves data one at a time
- ↳ aka low level

- ↳ uses 3rd party languages

Typical DBMS Component Modules

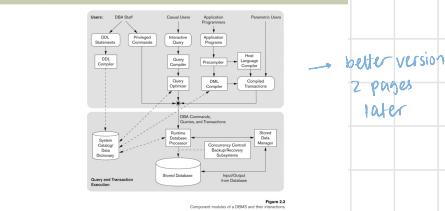


Figure 2.2 Component modules of a DBMS and their connections

Slide 2-21

Structural → SQL



→ no SQL

- ↳ relates to DB schema and organisation
- ↳ create, modify, delete objects → tables, indexes, views, constraints
- ↳ defines structure, relationships and constraints
- ↳ changes can impact entire DB system

no of students
course name

↳ google

? ↗ non Structural

- ↳ focuses on actual content/records in DB
↳ records specific info into
↳ documents in structure
- ↳ querying, inserting, updating, deleting data in tables
- ↳ manipulates info
- ↳ more common
- ↳ less likely to affect overall DB design

- ↳ mostly used in Big data
if data is Variety Velocity volume
then non structural

Big data apps ↗ insta
fb ↗ sc

Centralized DBMS

- ↳ Combines everything into a single system
- ↳ front end, backend, DB
- ↳ Prototype of mobile app

→ aka one tier architecture

Clients

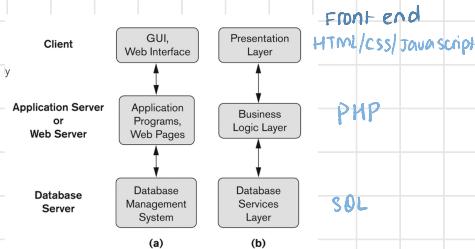
- ↳ Provide appropriate interfaces
- ↳ Connected to servers via network

DBMS Server

- ↳ Provide db query and transaction services to clients

Three-Tier Client - server

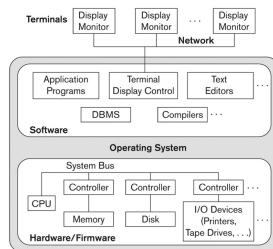
- ↳ websites
- ↳ no direct access to server
- ↳ enhanced security



Three Tier Client-Server Architecture

- Common for Web applications
- Intermediate Layer called Application Server or Web Server:
 - Stores the web connectivity software and the business logic part of the application used to access the corresponding data from the database server
 - Acts like a conduit for sending partially processed data between the database server and the client
- Three-tier Architecture Can Enhance Security:
 - Database server only accessible via middle tier
 - Clients cannot directly access database server
 - Clients contain user interfaces and Web browsers
 - The client is typically a PC or a mobile device connected to the Web

A Physical Centralized Architecture



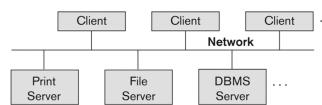
Copyright © 2016 Ramez Elmasri and Shamkant B. Navathe

Figure 2.4
A physical centralized architecture.

Slide

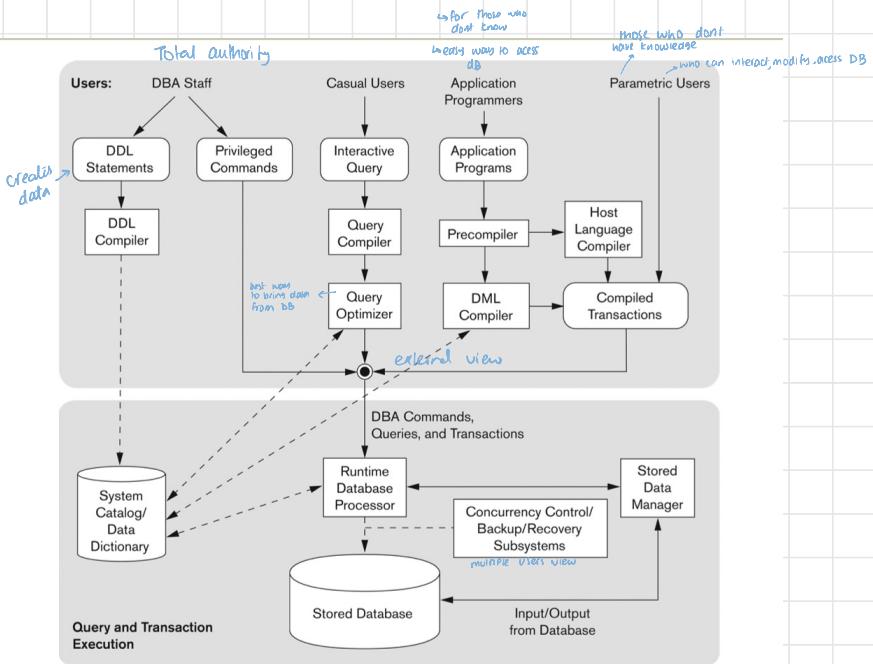
Two-Tier Client - server

- ↳ for organisational
- ↳ separate client and server
- ↳ client direct access to server → disaster



Two Tier Client-Server Architecture

- Client and server must install appropriate client module and server module software for ODBC or JDBC
- A client program may connect to several DBMSs, sometimes called the data sources.
- In general, data sources can be files or other non-DBMS software that manages data.
- See Chapter 10 for details on Database Programming



RELATIONAL MODEL CONCEPTS

↓ based on

Relation

- ↳ mathematical concept based on sets
- ↳ contains a set of rows *→ tuples*

STUDENT						
Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	NULL	19	3.21
Chung-cha Kim	381-62-1245	375-4409	125 Kirby Road	NULL	18	2.89
Dick Davidson	422-11-2320	NULL	3452 Elgin Road	749-1253	25	3.53
Rohan Panchal	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	NULL	19	3.25

Figure 5.1
The attributes and tuples of a relation STUDENT.

Formal Definitions - Tuple

- A **tuple** is an ordered set of values (enclosed in angled brackets < ... >)
- Each value is derived from an appropriate **domain**.
- A row in the CUSTOMER relation is a 4-tuple and would consist of four values, for example:

 - <632895, "John Smith", "101 Main St. Atlanta, GA 30332", "(404) 894-2000">
 - This is called a 4-tuple as it has 4 values
 - A tuple (row) in the CUSTOMER relation.

- A relation is a set of such tuples (rows)

Schema of Relation

CUSTOMER (Cust-id, Cust-name, Address, Phone#)

- CUSTOMER is the **relation name**
- Defined over the four attributes: Cust-id, Cust-name, Address, Phone#
- Each attribute has a **domain** or a set of valid values.
 - For example, the domain of Cust-id is 6 digit numbers.

Definition Summary

Informal Terms	Formal Terms
Table	Relation
Column Header	Attribute
All possible Column Values	Domain
Row	Tuple
Table Definition	Schema of a Relation
Populated Table	State of the Relation

- All values are considered atomic (indivisible).
- Each value in a tuple must be from the domain of the attribute for that column

CONSTRAINTS

Constraints determine which values are permissible and which are not in the database.

They are of three main types:

1. **Inherent or Implicit Constraints:** These are based on the data model itself. (E.g., relational model does not allow a list as a value for any attribute)
2. **Schema-based or Explicit Constraints:** They are expressed in the schema by using the facilities provided by the model. (E.g., max. cardinality/ratio constraint in the ER model)
3. **Application based or semantic constraints:** These are beyond the expressive power of the model and must be specified and enforced by the application programs.

- There are three *main types* of (explicit schema-based) constraints that can be expressed in the relational model:
 - **Key constraints**
 - **Entity integrity constraints**
 - **Referential integrity constraints**
- Another schema-based constraint is the **domain** constraint
 - Every value in a tuple must be from the *domain of its attribute* (or it could be **null**, if allowed for that attribute)

Relational Database State

- A **relational database state** DB of S is a set of relation states DB = { r_1, r_2, \dots, r_m } such that each r_i is a state of R_i and such that the r_i relation states satisfy the integrity constraints specified in IC.
- A relational database state is sometimes called a relational database *snapshot* or *instance*.
- We will not use the term *instance* since it also applies to single tuples.
- A database state that does not meet the constraints is an invalid state

Populated database state

- Each *relation* will have many tuples in its current relation state
- The *relational database state* is a union of all the individual relation states
- Whenever the database is changed, a new state arises
- Basic operations for changing the database:
 - **INSERT** a new tuple in a relation
 - **DELETE** an existing tuple from a relation
 - **MODIFY** an attribute of an existing tuple
- Next slide (Fig. 5.6) shows an example state for the COMPANY database schema shown in Fig. 5.5.

Entity Integrity

↳ More than 1 table

↳ Primary key unique can't be null

↳ Foreign key can't be null
name can be different
domain will be that of primary key

be duplicated
be null

Delete

if deleted all other tables will

↳ set null be null

↳ set default be default value set

↳ cascade like nested loops
would make an loops
null when delete

↳ reject

↳ restrict

EMPLOYEE

E.no	E.id	F.name	CNIC	Gender	Ph no	Addl	D.no
------	------	--------	------	--------	-------	------	------

Primary key

don't create in Employee table
add as foreign key

Department

D.no	Dname	D.CNIC	Manager	D.start date	D.location
------	-------	--------	---------	--------------	------------

Primary key

NULL

123 → resigning

SD will be deleted and

Will be null in other duplicate tables

Location

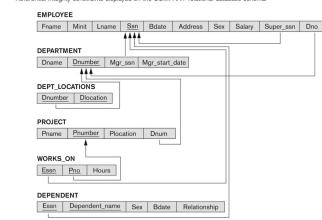
D.no	D.location
------	------------

Foreign key part of primary key so can't null but can be duplicated

combine both
and make the whole primary key

Rerential Integrity Constraints for COMPANY database

Figure 5.2
Referential integrity constraints displayed on the COMPANY relational database schema.



Stored procedures

↳ functions that are stored already

and called when needed

FALVIN

concatenate columns

column1 || column2 As joined

violated by each operation, if any, and the different ways of enforcing these constraints.

- Insert <'Robert', 'F', 'Scott', '943775543', '1972-06-21', '2365 Newcastle Rd, Bellaire, TX', M, 58000, 888665555, 1> into EMPLOYEE. NO VIOLATION
 - Insert <'ProductA', 4, 'Bellaire', 2> into PROJECT. REFERENTIAL constraint
RE: ProductA
RE: Bellaire
 - Insert <'Production', 4, '943775543', '2007-10-01'> into DEPARTMENT. KEY constraint
RE: Production
RE: 943775543
 - Insert <677678989, NULL, '40.0' into WORKS_ON. no constraint
no value
 - Insert <'453453453', 'John', 'M', '1990-12-12', 'spouse'> into DEPENDENT. no violation
 - Delete the WORKS_ON tuples with Essn = '333445555'. no violation
 - Delete the EMPLOYEE tuple with Ssn = '987654321'. REFERENTIAL constraint
RE: John
RE: 987654321
 - Delete the PROJECT tuple with Pname = 'ProductX'. no constraint
RE: ProductX
RE: 987654321
 - Modify the Mgr_ssn and Mgr_start_date of the DEPARTMENT tuple with Dnumbr = 5 to '123456789' and '2007-10-01', respectively. no constraint
 - Modify the Super_ssn attribute of the EMPLOYEE tuple with Ssn = '999887777' to '943775543'. no constraint
RE: Super_ssn
RE: 999887777
RE: 943775543
 - Modify the Hours attribute of the WORKS_ON tuple with Essn = '999887777' and Pno = to '5.0'. domain constraint
RE: Hours
RE: 999887777
RE: 5.0
- 5.12. Consider the AIRLINE relational database schema shown in Figure 5.8, which describes a database for airline flight information. Each FLIGHT is identified by a Flight_number, and consists of one or more FLIGHT_LEGs with Leg_numbers 1, 2, 3, and so on. Each FLIGHT_LEG has scheduled arrival and departure times, airports, and one or more LEG_INSTANCES—one for each Date on which the flight travels. FAREs are kept for each FLIGHT. For each FLIGHT_LEG instance, SEAT_RESERVATIONS are kept, as are the AIRPLANE used on the leg and the actual arrival and departure times and airports. An AIRPLANE is identified by an Airplane_id and is of a particular AIRPLANE_TYPE. CAN_LAND relates AIRPLANE_TYPES to the AIRPORTS at which they can land. An AIRPORT is identified by an Airport_code. Consider an update for the AIRLINE database to enter a reservation on a particular flight or flight leg on a given date.
- Give the operations for this update.
 - What types of constraints would you expect to check?
 - Which of these constraints are key, entity integrity, and referential integrity constraints, and which are not?
 - Specify all the referential integrity constraints that hold on the schema shown in Figure 5.8.

- 5.13. Consider the relation CLASS(Course#, Univ_Section#, Instructor_name, Semester, Building_code, Room#, Time_period, Weekdays, Credit_hours). This represents classes taught in a university, with unique Univ_section#. Identify what you think should be various candidate keys, and write in your own words the conditions or assumptions under which each candidate key would be valid.

INSERT may violate any of the constraints:

- Domain constraint: attribute type wrong → update
 - if one of the attribute values provided for the new tuple is not of the specified attribute domain
- Key constraint: → same value of pk
 - if the value of a key attribute in the new tuple already exists in another tuple in the relation
- Referential integrity: → related to fk
 - if a foreign key value in the new tuple references a primary key value that does not exist in the referenced relation
- Entity integrity: → pk is null
 - if the primary key value is null in the new tuple

Figure 5.6
One possible database state for the COMPANY relational database schema.

The diagram illustrates the COMPANY relational database schema with five tables:

- EMPLOYEE** (ER diagram: `Employee`)
- DEPARTMENT** (ER diagram: `Department`)
- WORKS_ON** (ER diagram: `WorksOn`)
- PROJECT** (ER diagram: `Project`)
- DEPENDENT** (ER diagram: `Dependent`)

Relationships shown:

- EMPLOYEE** → **DEPARTMENT**: M:N (many-to-many)
- EMPLOYEE** → **WORKS_ON**: M:N (many-to-many)
- DEPARTMENT** → **PROJECT**: M:N (many-to-many)
- WORKS_ON** → **PROJECT**: M:N (many-to-many)
- DEPENDENT** → **EMPLOYEE**: 1:M (one-to-many)

EMPLOYEE table data:

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1995-12-08	638 Voss, Houston, TX	M	40000	333445555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	333445555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabber	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	NULL

DEPARTMENT table data:

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

WORKS_ON table data:

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT table data:

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT table data:

Essn	Dependent_name	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1989-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Almer	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

→ FK not in PK
delete PK, which is also a FK somewhere
FK domain

Possible violations for each operation

- **INSERT** may violate any of the constraints:
 - Domain constraint:
 - if one of the attribute values provided for the new tuple is not of the specified attribute domain
 - Key constraint:
 - if the value of a key attribute in the new tuple already exists in another tuple in the relation
 - Referential integrity:
 - if a foreign key value in the new tuple references a primary key value that does not exist in the referenced relation
 - Entity integrity:
 - if the primary key value is null in the new tuple
- **DELETE** may violate only referential integrity:
 - If the primary key value of the tuple being deleted is referenced from other tuples in the database
 - Can be remedied by several actions: RESTRICT, CASCADE, SET NULL (see Chapter 6 for more details)
 - RESTRICT option: reject the deletion
 - CASCADE option: propagate the new primary key value into the foreign keys of the referencing tuples
 - SET NULL option: set the foreign keys of the referencing tuples to NULL
 - One of the above options must be specified during database design for each foreign key constraint
- **UPDATE** may violate domain constraint and NOT NULL constraint on an attribute being modified
- Any of the other constraints may also be violated, depending on the attribute being updated:
 - Updating the primary key (PK):
 - Similar to a DELETE followed by an INSERT
 - Need to specify similar options to DELETE
 - Updating a foreign key (FK):
 - May violate referential integrity
 - Updating an ordinary attribute (neither PK nor FK):
 - Can only violate domain constraints

SPECIFYING CONSTRAINTS IN SQL

Basic constraints:

- Relational Model has 3 basic constraint types that are supported in SQL:
 - **Key constraint:** A primary key value cannot be duplicated
 - **Entity Integrity Constraint:** A primary key value cannot be null
 - **Referential integrity** constraints : The “foreign key” must have a value that is already present as a primary key, or may be null.

SQL CREATE TABLE data definition statements for defining the COMPANY schema from Figure 5.7 (Fig. 6.1)

```

CREATE TABLE EMPLOYEE
  (Fname          VARCHAR(15),
   Minit          CHAR,
   Lname          VARCHAR(15),
   Ssn            CHAR(9),
   Bdate          DATE,
   Address        VARCHAR(30),
   Sex            CHAR,
   Salary          DECIMAL(10,2),
   Super_ssn      CHAR(9),
   Dno            INT)
   PRIMARY KEY (Ssn),
CREATE TABLE DEPARTMENT
  (Dname          VARCHAR(15),
   Dnumber         INT,
   Manager        CHAR(9),
   Mgr_start_date DATE)
   PRIMARY KEY (Dnumber),
   UNIQUE (Dname),
   FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn);
CREATE TABLE DEPT_LOCATIONS
  (Dnumber         INT,
   Location        VARCHAR(15))
   PRIMARY KEY (Dnumber, Location),
   FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber);

```

continued on next slide

Attribute Data Types and Domains in SQL

▪ Basic data types

▪ Numeric data types

- Integer numbers: INTEGER, INT, and SMALLINT
- Floating-point (real) numbers: FLOAT or REAL, and DOUBLE PRECISION

▪ Character-string data types

- Fixed length: CHAR(n), CHARACTER(n)
- Varying length: VARCHAR(n), CHAR VARYING(n), CHARACTER VARYING(n)

▪ Bit-string data types

- Fixed length: BIT(n)
- Varying length: BIT VARYING(n)

▪ Boolean data type

- Values of TRUE or FALSE or NULL

▪ DATE data type

- Ten positions
- Components are YEAR, MONTH, and DAY in the form YYYY-MM-DD
- Multiple mapping functions available in RDBMSs to change date formats

▪ Additional data types

▪ Timestamp data type

Includes the DATE and TIME fields

- Plus a minimum of six positions for decimal fractions of seconds
- Optional WITH TIME ZONE qualifier

▪ INTERVAL data type

- Specifies a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp
- DATE, TIME, Timestamp, INTERVAL data types can be cast or converted to string formats for comparison.

▪ PRIMARY KEY clause

- Specifies one or more attributes that make up the primary key of a relation
- Dnumber INT PRIMARY KEY;

▪ UNIQUE clause

- Specifies alternate (secondary) keys (called CANDIDATE keys in the relational model).
- Dname VARCHAR(15) UNIQUE;

Specifying Constraints on Tuples Using CHECK

- Additional Constraints on individual tuples within a relation are also possible using CHECK
- CHECK clauses at the end of a CREATE TABLE statement
- Apply to each tuple individually
- CHECK (Dept_create_date <= Mgr_start_date);

Arithmetic Operators

▪ LIKE comparison operator

- Used for string pattern matching
- % replaces an arbitrary number of zero or more characters
- underscore (_) replaces a single character
- Examples: WHERE Address LIKE '%Houston,TX%';
- WHERE Ssn LIKE '_ _ 1 _ 8901';

▪ BETWEEN comparison operator

E.g., in Q14 :

```
WHERE(Salary BETWEEN 30000 AND 40000)
      AND Dno = 5;
```

▪ Aliases or tuple variables

- Declare alternative relation names E and S to refer to the EMPLOYEE relation twice in a query:

Query 8. For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

```
▪ SELECT E.Fname, E.lname, S.fname, S.lname
  FROM   EMPLOYEE AS E, EMPLOYEE AS S
  WHERE E.Super_ssn=S.Ssn; /*using 's' will make join
```

- Recommended practice to abbreviate names and to prefix same or similar attribute from multiple tables.

U1: **INSERT INTO** EMPLOYEE
VALUES ('Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
Oak Forest, Katy, TX', 'M', 37000, '653298653', 4);

The variation below inserts multiple tuples ...

BULK LOADING OF TABLES

- Another variation of **INSERT** is used for bulk-loading of several tuples into tables
- A new table TNEW can be created with the same attributes as T and using LIKE and DATA in the syntax, it can be loaded with entire data.
- EXAMPLE:

```
CREATE TABLE DSEMPs LIKE EMPLOYEE  
(SELECT E.*  
  FROM   EMPLOYEE AS E  
 WHERE   E.Dno=5)  
WITH DATA;
```

- Example: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively

U5: **UPDATE** PROJECT
SET PLOCATION = 'Bellaire',
 DNUM = 5
WHERE PNUMBER=10

DELETE

- Removes tuples from a relation
 - Includes a WHERE-clause to select the tuples to be deleted
 - Referential integrity should be enforced
 - Tuples are deleted from only *one table* at a time (unless CASCADE is specified on a referential integrity constraint)
 - A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table
 - The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause

U4A: **DELETE FROM** EMPLOYEE
WHERE Lname='Brown';

foreign key can't be null

Primary key can't be null

IN \Rightarrow include

* **Select** → columns

* **From** → Tables

* **Where** → rows, tuples, records

→ conditional

* **Order by** → DESC, ASC

repetition

* **Group by** → count, min, max, avg, sum

↳ always with select clause

↳ count never counts 'NULL'

↳ group by without functions give distinct clause

Select

count (DISTINCT(column-name/*))

As count_no

From Table

SELECT

model_year,

(count (product_id) AS product_count)

FROM Products

GROUP BY model_year, as repeated

→ can't write orderby here,
as aggregate function

SELECT

count(CASE WHEN price >= 500 THEN product_id END)

AS expensive

FROM Products

Select Dep

From Departments

Group by Dep

↓
only one column
can hold more UNLESS
others are aggregate

* **Having** → group by

→ conditional

↳ doesn't work on single clause

↳ on complete group

* **Aggregate function**

↳ always use group by with it

↳ can't use distinct with aggregate function

↳ count, min, max, avg, sum

EMPLOYEE

E_id	E_name	D no
123	Ali	1
321	Bilal	3
419	Ahmed	3
512	Hassan	1

Department

D_no	Dname	Dmang_id
1	CS	512
2	EE	NULL
3	AI	321

Select * from EMPLOYEE E, Department D where E.E_id = D.Dmang_id

E_id	E-name	Dno	Dno	Dname	Dmang_id
123	Ali	1	1	CS	512
123	Ali	1	2	EE	NULL
123	Ali	1	3	AI	321
321	Bilal	3	1	CS	512
321	Bilal	3	2	EE	NULL
321	Bilal	3	3	AI	321
419	Ahmed	3	1	CS	512
419	Ahmed	3	2	EE	NULL
419	Ahmed	3	3	AI	321
512	Hassan	1	1	CS	512
512	Hassan	1	2	EE	NULL
512	Hassan	1	3	AI	321

→ CROSS PRODUCT

* if there's a common column
it should be a primary key

→ might not work on
oracle compiler as it
can't support

Select * from EMPLOYEE E, Department D where E.E_id = D.Dmang_id

E_id	E-name	Dno	Dno	Dname	Dmang_id
321	Bilal	3	3	AI	321
512	Hassan	1	1	CS	512

Select * from EMPLOYEE E, Department D where E.E_id=D.Dmang_id AND D.Dname='CS'

E_id	E-name	Dno	Dno	Dname	Dmang_id
512	Hassan	1	1	CS	512

by default inner join

JOIN

↳ necessary b/w common columns

better query optimization compared to SQL

→ b/w two tables

explain/analyse select * from Employee
create index on Employee using Btree (department)

CROSS JOIN

→ diff tables
combiner product

↳ doesn't have where clause

Select * from Employee, Department

→ join clause

SELF JOIN

→ same tables

Select E.fname, E.lname, S.fname, S.lname

from Employee E JOIN Employee S

ON E.ssn = S.superssn

→ self and inner join

INNER JOIN

↳ returns matching rows

NATURAL JOIN

If Table1 and Table2 have same column name

Select * from Department D

JOIN dept-locations

USING D.number

EQUI JOIN

if ' = ' written

Select E.fname, E.lname, S.fname, S.lname

from Employee E, Employee S

where E.ssn = S.superssn

\downarrow
if not written
cross join

OUTER JOIN

↳ returns matching rows

↳ to check where NULL comes

LEFT JOIN

↳ full info of left table and some info of right

LEFT OUTER JOIN

FULL JOIN

↳ full info of left and right table

FULL JOIN

RIGHT JOIN

↳ full info of right table and some info of left

↳ if no matching rows then displays NULL

HOW TO JOIN more than 2 Tables

Select column.name

\downarrow
not in Paper

From T1 Join T2

ON column.name

JOIN T3

ON T2._ = T3._

RIGHT JOIN

EQUAL JOIN

- very expensive
- too many computations
- slows system down

which is better

Sub Query OR JOIN

- ↳ if no common column
- ↳ if no of tables less
- ↳ easier to read
- ↳ filter or aggregate data

- ↳ on common columns
- ↳ if no of tables greater
- ↳ retrieve data from multiple tables simultaneously

```
START_DATE = TO_DATE('2001-01-13', 'YYYY-MM-DD'));
```

TO_CHAR (hire_date, 'Month')
 Extract (year From hire_date)
 TO_DATE (hire_date, 'DD-MM-YYYY')

name like '%a_'
 substr (name, -3)
 round (num, 1)

Conclusion for Joins:

SELECT <fields>
 FROM TableA A
 INNER JOIN TableB B
 ON A.key = B.key

SELECT <fields>
 FROM TableA A
 RIGHT JOIN TableB B
 ON A.key = B.key

SELECT <fields>
 FROM TableA A
 LEFT JOIN TableB B
 ON A.key = B.key
 WHERE B.key IS NULL

**SQL
JOINS**

SELECT <fields>
 FROM TableA A
 RIGHT JOIN TableB B
 ON A.key = B.key
 WHERE A.key IS NULL

SELECT <fields>
 FROM TableA A
 FULL OUTER JOIN TableB B
 ON A.key = B.key

SELECT <fields>
 FROM TableA A
 FULL OUTER JOIN TableB B
 ON A.key = B.key
 WHERE A.key IS NULL
 OR B.key IS NULL

ishma-hafeez
notes
reprst
refct

This work is licensed under a Creative Commons Attribution 3.0 Unported License.
 Author: <http://commons.wikimedia.org/wiki/User:Arbeck>

INTRODUCTION TO SET OPERATOR

Query 18. Retrieve the names of all employees who do not have supervisors.

```
Q18: SELECT Fname, Lname
      FROM EMPLOYEE
      WHERE Super_ssn IS NULL;
```

Nested Queries (cont'd.)

- Use tuples of values in comparisons
 - Place them within parentheses

```
SELECT DISTINCT Essn
  FROM WORKS_ON
 WHERE (Pno, Hours) IN ( SELECT Pno, Hours
      FROM WORKS_ON
     WHERE Essn='123456789');
```

- Use other comparison operators to compare a single value v
 - = ANY (or = SOME) operator
 - Returns TRUE if the value v is equal to some value in the set V and is hence equivalent to IN
 - Other operators that can be combined with ANY (or SOME): >, >=, <, <=, and <>
 - ALL: value must exceed all values from nested query


```
SELECT Lname, Fname
    FROM EMPLOYEE
   WHERE Salary > ALL ( SELECT Salary
      FROM EMPLOYEE
     WHERE Dno=5 );
```

Copyright © 2016 Ramez Elmasri and Shamkant B. Navathe

Slide 7- 11

- Queries that are nested using the = or IN comparison operator can be collapsed into one single block: E.g., Q16 can be written as:
- Q16A:


```
SELECT E.Fname, E.Lname
        FROM EMPLOYEE AS E, DEPENDENT AS D
       WHERE E.Ssn=D.Essn AND E.Sex=D.Sex
             AND E.Fname=D.Dependent_name;
```
- Correlated nested query
 - Evaluated once for each tuple in the outer query

USE of EXISTS

Q7:

```
SELECT Fname, Lname
  FROM Employee
 WHERE EXISTS (SELECT *
      FROM DEPENDENT
     WHERE Ssn= Essn)

 AND EXISTS (SELECT *
      FROM Department
     WHERE Ssn= Mgr_Ssn)
```

Query: List first and last name of employees who work on **ALL** projects controlled by Dno=5.

```
SELECT Fname, Lname
  FROM Employee
 WHERE NOT EXISTS ( (SELECT Pnumbr
      FROM PROJECT
     WHERE Dno=5)

 EXCEPT (SELECT Pno
      FROM WORKS_ON
     WHERE Ssn= Essn))
```

The above is equivalent to double negation: List names of those employees for whom there does NOT exist a project managed by department no. 5 that they DO NOT work on.

Copyright © 2016 Ramez Elmasri and Shamkant B. Navathe

Slide 7- 16

- Can use explicit set of values in WHERE clause

```
Q17: SELECT DISTINCT Essn
      FROM WORKS_ON
     WHERE Pno IN (1, 2, 3);
```

Example: LEFT OUTER JOIN

```
SELECT E.Lname AS Employee_Name
      , S.Lname AS Supervisor_Name
  FROM Employee AS E LEFT OUTER JOIN EMPLOYEE AS S
    ON E.Super_ssn = S.Ssn
```

ALTERNATE SYNTAX:

```
SELECT E.Lname , S.Lname
  FROM EMPLOYEE E, EMPLOYEE S
 WHERE E.Super_ssn + = S.Ssn
```

Renaming Results of Aggregation

- Following query returns a single row of computed values from EMPLOYEE table:

```
Q19:   SELECT SUM (Salary), MAX (Salary), MIN (Salary), AVG
                  (Salary)
        FROM EMPLOYEE;
```

- Specifies grouping attributes
- COUNT (*) counts the number of rows in the group

Grouping: The GROUP BY and HAVING Clauses (cont'd.)

■ HAVING clause

- Provides a condition to select or reject an entire group:
- Query 26.** For each project on which more than two employees work, retrieve the project number, the project name, and the number of employees who work on the project.

```
Q26:      SELECT      Pnumber, Pname, COUNT (*)
          FROM        PROJECT, WORKS_ON
          WHERE       Pnumber=Pno
          GROUP BY    Pnumber, Pname
          HAVING      COUNT (*) > 2;
```

Combining the WHERE and the HAVING Clause

- Consider the query: we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.

■ INCORRECT QUERY:

```
SELECT      Dno, COUNT (*)
FROM        EMPLOYEE
WHERE       Salary>40000
GROUP BY    Dno
HAVING      COUNT (*) > 5;
```

Correct Specification of the Query:

- Note: the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

Query 28. For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

```
Q28:     SELECT      Dnumber, COUNT (*)
          FROM        DEPARTMENT
          WHERE       Dnumber=Dno AND Salary>40000 AND
                     ( SELECT    Dno
                       FROM      EMPLOYEE
                       GROUP BY Dno
                       HAVING    COUNT (*) > 5 )
```

Alter Command

- Add Column/s
- Remove Column/s
- Modify datatype
- Modify datatype length
- Add Constraints
- Remove Constraints
- Rename column / table

```
alter table employee add address varchar(10);
desc emp1
alter table employee drop column address;
alter table employee modify id varchar(30);
alter table employee rename column id to roll_no;
alter table employee rename to emp1;
```

Alter

DDL

- ✓
- ✓
- ✓
- ✓ int
- ✓ ↓
- ✓ varchar

Name varchar(10)

ID → Eid

Emp → Emp_detail

Update

DML

Data.

Update Emp
Set Salary = Salary * 2
Where id=1;

Emp			
ID	Name	Salary	Email
1	A	10000	20000
2	B	20000	40000
3	C	30000	60000



SUBSCRIBED



1. capslock matters in 'strings'

2. Use IN for listing `IN ('101','102')` instead of :

→ Write a SQL query to display Second highest Salary from Emp table?

E_id	E_name	Dept	Salary
1	Ram	HR	10000
2	Amit	MRKT	20000
3	Ravi	HR	30000
4	Nitin	MRKT	40000
5	Varma	IT	50000

→ Write a SQL query to display Employee name who is having second highest salary?

Max
Select (Salary) from
Emp where Salary > (Select Max(Salary) from Emp)
Select (Salary) from
Emp where Salary < (Select Max(Salary) from Emp)
Select (Salary) from
Emp where Salary < (Select Max(Salary) from Emp)
Select (Salary) from
Emp where Salary < (Select Max(Salary) from Emp)

Select E_name from Emp where dept In
HR 2
MRKT 2
IT 1
dept from Emp group by dept
Having count(*) < 2

E_id	E_name	Dept	Salary
1	Ram	HR	10000
2	Amit	MRKT	20000
3	Ravi	HR	30000
4	Nitin	MRKT	40000
5	Varma	IT	50000

In / Not In
Using (query)
find the name of Emps who are not working on a project?

Nested → Bottom up

Find Ename from
Emp where E_id In
(Select E_id from Project)

Inner query
Project

PK	Eid	Ename	Address
1	Ram	Chd	
2	Varma	Delhi	
3	Nitin	Pune	
4	Robin	Banglore	
5	Ammy	Chd	

PK	Eid	Pid	Prname	Location
1	P1	IT	Banglore	
5	P2	Big Data	Delhi	
3	P3	Retail	Pune	
4	P4	Android	Mumbai	
2	P5	Test	Hyderabad	

In / Not In
Detail of Emp whose address is either Delhi or Chd or Pune;

Selected * from Emp where Address In ('Delhi', 'Chd', 'Pune')

Address = ('Delhi', 'Chd', 'Pune')
WRONG

Address = 'Delhi' and Address='Chd'
CORRECT

PK	Eid	Pid	Prname
1	P1	IT	
5	P2	Big Data	
3	P3	Retail	
4	P4	Android	

'Correlated Subquery (Synchronized query)'

→ It is a subquery that uses values from outer query.
→ Top down approach

Find all employees detail who works in a department.

Select * from Emp where dept in (Select dept from Dept where Dept_id = Emp_id);

Exists (Select * from Dept where Dept_id = Emp_id);

E_id	E_name	Dept	Sal
1	Ram	HR	10000
2	Amit	MRKT	20000
3	Ravi	HR	30000
4	Nitin	MRKT	40000
5	Varma	IT	50000

D_id	Dept	F_id
1	HR	1
2	MRKT	2
3	IT	3
4	Test	4

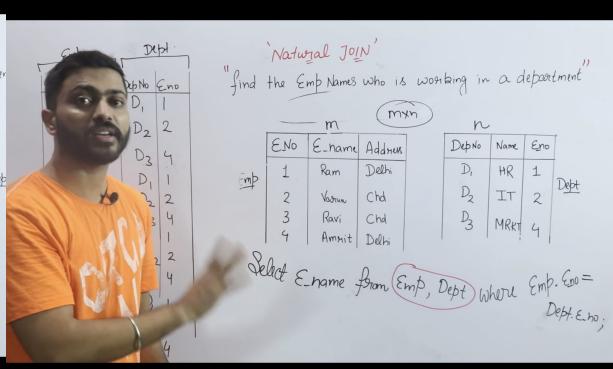
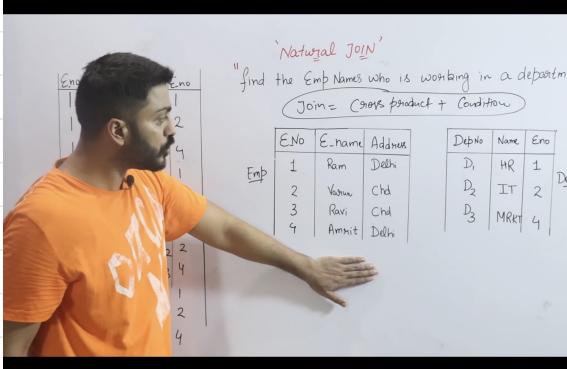
Nested Subquery
→ Bottom up
From inner query
Then outer query

Top down approach
From outer query
Then inner query

Join
→ Cross product + Condition
inner side from inner
then check condition

Outer query
Select * from Emp where exists (Select id from Dept where emp_id = dept_id);

Inner query
Select attribute from Emp, Dept where emp_id = dept_id;



Left Outer Join

it gives the matching rows and the rows which are in left table but not in right table.

Empno	E_name	d_name	Loc
E ₁	Karen	IT	Dell, I
E ₂	Amit	HR	Hd
E ₃	Ravi	IT	Dell, I
E ₄	Nitin		

Emp.no		Empl.name	Deptno	Dept	
				Dept.-No	D.name
E ₁	Vasan	D ₁		D ₁	IT
E ₂	Ammrit	D ₂		D ₂	HR
E ₃	Ravi	D ₁		D ₃	Finance
E ₄	Nitin	—			Print

Select emp.no, e-name, d-name, loc from emp left outer join dept
On (emp.dept.no = dept.dept.no)

```
SELECT * FROM jobs where job_title in ('President','Accountant');
```

- (d) **Pattern Match Search Condition:** The search condition involves searching for a particular character or string within a column value. Like Operator with the help of pattern matching symbols (_, %) are used find patterns in the column's value. '_' represents a single character while '%' represents a sequence of characters.

Example: List all the employees whose names contains an ‘a’ in their first names

```
SELECT * FROM employees WHERE first_name LIKE '%a%';
```

OR List all employees having L as second letter in their first names.

```
SELECT * FROM employees WHERE first_name LIKE ' a %';
```

```
bb3
Worksheet : Query Builder
25: --/
30: --_ALTER TABLE JOBS ADD PRIMARY KEY (JOB_ID);
31: --_ALTER TABLE JOBS MODIFY JOB_ID NUMBER;
32: --_ALTER TABLE JOBS_HISTORY ADD FOREIGN KEY (JOB_ID) REFERENCES JOBS(JOB_ID);
33: --_ALTER TABLE JOB_HISTORY ADD FOREIGN KEY (JOB_ID) REFERENCES JOBS(JOB_ID);
34: --
35: --Insert into JOBS VALUES(4468, "manager", 3400, 6544);
36: --SELECT * FROM JOBS;
37: --_ALTER TABLE JOBS ADD Job_Nature VARCHAR2(35);
38: --
39: --_ALTER TABLE EMPLOYEES ADD primary key (EMPLOYEE_ID);
40: ALTER TABLE JOB_HISTORY ADD FOREIGN KEY (EMPLOYEE_ID) REFERENCES EMPLOYEES(EMPLOYEE_ID);
41: --
```

```
56  
57 --UPDATE EMPLOYEES SET SALARY = 8000 WHERE EMPLOYEE_ID = '105' AND SALARY <1000;  
58  
59 ALTER TABLE EMPLOYEES ADD CONSTRAINT PK_Employee_Job PRIMARY KEY (EMPLOYEE_ID, JOB_ID);  
60  
61 --DROP TABLE EMPLOYEES;  
62  
63
```

- **COMMENT** - add comments to the data dictionary
- **RENAME** - rename an object

1. Create

There are two CREATE statements available in SQL:

- CREATE DATABASE

Example Syntax:

- **CREATE DATABASE** database_name;
database_name: name of the database.

- CREATE TABLE

Example Syntax:

- **CREATE TABLE** table_name {column1 data_type(size), column2 data_type(size), column3 data_type(size)...};
table_name: name of the table.
column1 name of the first column.
data_type: Type of data we want to store in the particular column.
size: Size of the data we can store in a particular column.

2. Alter

ALTER TABLE is used to add, delete/drop or modify columns in the existing table. It is also used to add and drop various constraints on the existing table.

Example Syntax:

ADD

- **ALTER TABLE** table_name **ADD** (Columnname_1 datatype, Columnname_2 datatype, ... Columnname_n datatype);

Drop

- **ALTER TABLE** table_name **DROP COLUMN** column_name;

Modify

- **ALTER TABLE** table_name **MODIFY** column_name column_type;

QUESTION 4

1. Define a table constraint on Emp that will ensure that every employee makes at least \$10,000.

```
ALTER TABLE Emp
ADD CONSTRAINT MinSalaryCheck CHECK (salary >= 10000);
```

2. Define a table constraint on Dept that will ensure that all managers have age > 30.

```
ALTER TABLE Dept
ADD CONSTRAINT ManagerAgeCheck CHECK (managerid IS NULL OR age > 30);
```

3. Drop

DROP is used to delete a whole database or just a table. The DROP statement destroys the objects like an existing database, table, index, or view. A DROP statement in SQL removes a component from a relational database management system [RDBMS].

Example Syntax:

- **DROP TABLE** table_name;
table_name: Name of the table to be deleted.
- **DROP DATABASE** database_name;
database_name: Name of the database to be deleted.