

Operation Systems

What do i not know?
No giving assignments! (except BU)
attendance doesn't matter much

CORNELL NOTEMAKING METHOD
WRONG

Project + presentation (10%)

Assignments (4 off) (10%)

Mid term ($15+15=30\%$)

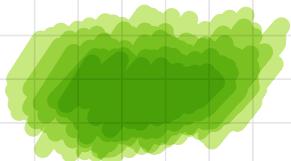
Final (50%)



ISHMA METHOD

POWER → SHELL

KILL
KILL



Virtualization:

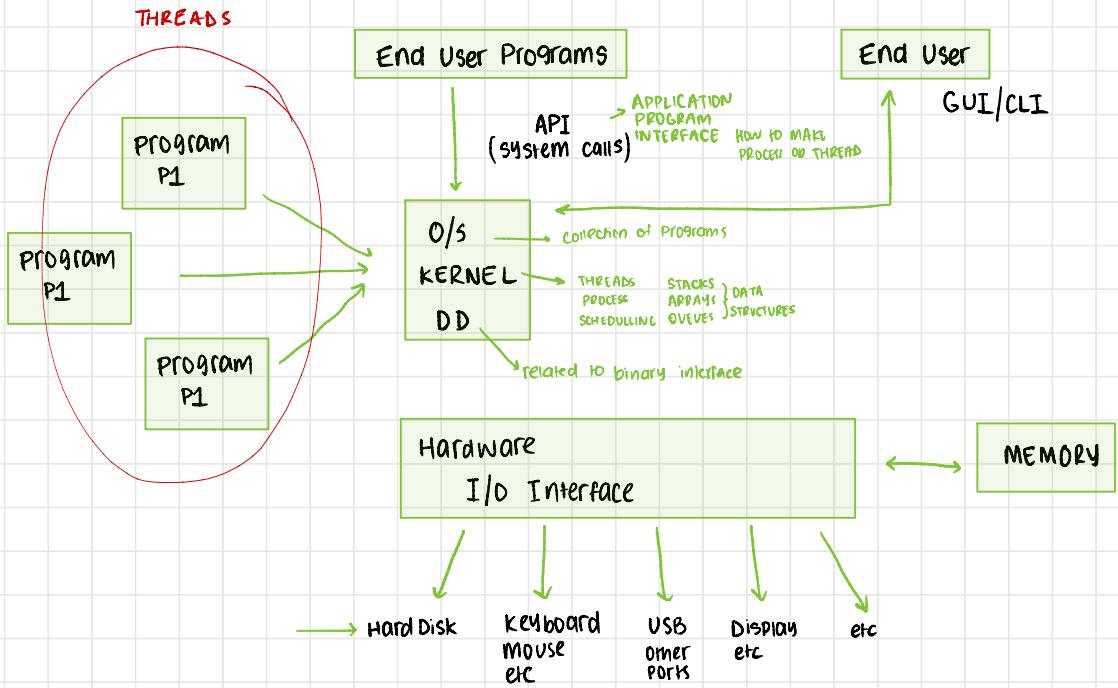
Let us take the most basic of resources, the CPU. Assume there is one physical CPU in a system (though now there are often two or four or more). What virtualization does is take that single CPU and make it look like many virtual CPUs to the applications running on the system. Thus, while each application thinks it has its own CPU to use, there is really only one. And thus the OS has created a beautiful illusion: it has virtualized the CPU.

Concurrency: There are certain types of programs that we call multi-threaded applications; each thread is kind of like an independent agent running around in this program, doing things on the program's behalf. But these threads access memory. If we don't coordinate access to memory between threads, the program won't work as expected. First, the OS must support multi-threaded applications with protection (such as locks and condition variables). Second, the OS itself was the first concurrent program — it must access its own memory very carefully or many strange and terrible things will happen.

Persistence: Making information survive computer crashes, disk failures, or power outages is a tough and interesting challenge and needs persistence.

Operating System

AB1
APPLICATION BINARY INTERFACE



3 CONCEPTS in O/S

architecture, H/W for each user programs

How to run more than one instruction stream in one program Using API's

- Virtualization → PROCESS IS A VIRTUALISATION → not in BATCH give access of all process to hardware together

- Concurrency

multiple threads → inside the process



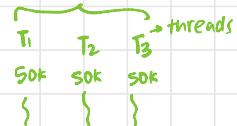
- Persistence

not in memory data/program storage that can survive reboots

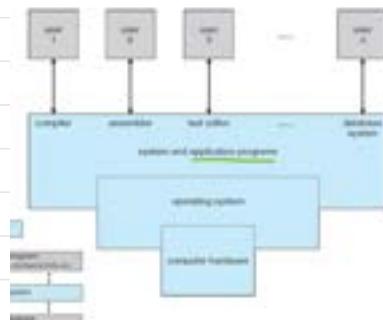
for (i=0 200000000

a[i] = b[i] + c[i]

divide data into chunks



- Computer system can be divided into four components:
 - **Hardware** - provides basic computing resources
 - **Operating system** - controls, monitors and coordinates integrated activities of computer, which includes application programs
 - **Application programs** - logic which is run according to user command requests
 - **User**



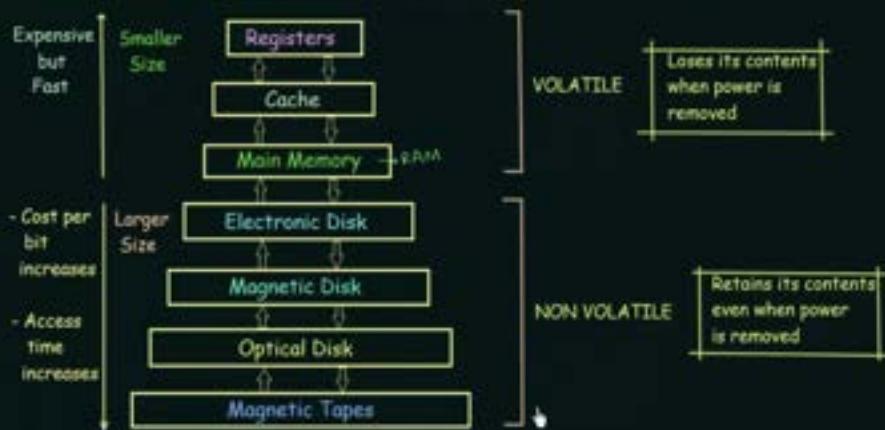
→ To ensure orderly access to the shared memory, a memory controller is provided whose function is to synchronize access to the memory

Some important terms:

- 1) **Bootstrap Program:** → The initial program that runs when a computer is powered up or rebooted.
 - It is stored in the ROM.
 - It must know how to load the OS and start executing that system.
 - It must locate and load into memory the OS Kernel.
- 2) **Interrupt:** → The occurrence of an event is usually signalled by an Interrupt from Hardware or Software.
 - Hardware may trigger an interrupt at any time by sending a signal to the CPU, usually by the way of the system bus.
- 3) **System Call (Monitor call):** → Software may trigger an interrupt by executing a special operation called System Call.

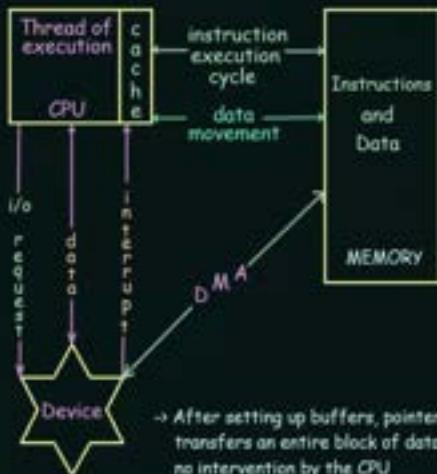
So, that Operating system is
already residing or stored

Basics of Operating System (Storage Structure)



which basically are secondary storage devices, are nonvolatile.

Working of an I/O Operation:



- To start an I/O operation, the device driver loads the appropriate registers within the device controller.
- The device controller, in turn, examines the contents of these registers to determine what action to take.
- The controller starts the transfer of data from the device to its local buffer.
- Once the transfer of data is complete, the device controller informs the device driver via an interrupt that it has finished its operation.
- The device driver then returns control to the operating system.

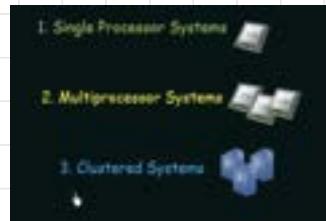
This form of interrupt-driven I/O is fine for moving small amounts of data but can produce high overhead when used for bulk data movement.

- To solve this problem, Direct Memory Access (DMA) is used
- After setting up buffers, pointers, and counters for the I/O device, the device controller transfers an entire block of data directly to or from its own buffer storage to memory, with no intervention by the CPU.

So, from these registers, the appropriate registers that are required for

Single Processor System

- ↳ single processor
- ↳ has multiple special purpose processor which handle specific tasks

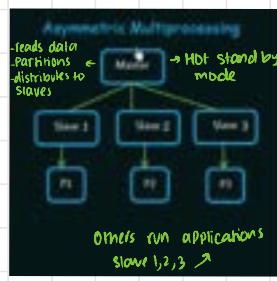
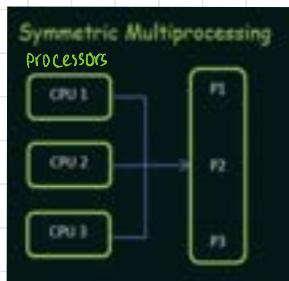


Multiprocessor System

- ↳ has 2 or more processors
- ↳ they share clock/memory/devices/bus

ADV

- ↳ increased throughput → amount of data transferred from 1 place to another → more processors
- ↳ economy of sales → as processors share resources
- ↳ increased reliability → if 1 processor fails other will work



- ↳ all CPUs similar
- ↳ all processors access same memory

- ↳ master and slave
- ↳ all processors access different memory
- ↳ independent memory

Clustered Systems

- ↳ multiple processors
- ↳ composed of 2 or more individual systems

ADV

- ↳ high availability → if 1 system fails the others can take care of task
- ↳ can also be structured symmetrically or asymmetrically

OS TYPES

if didn't exist then 2 job would take all CPU time until completion in which vs 2 I/O devices CPU would be idle

Multiprogramming

- ↳ executes multiple jobs
- ↳ no user interaction with comp system
- ↳ processes are in RAM
- ↳ only one process executed at a time
- ↳ if I/O operation occurs then other process executed

Time sharing / Multitasking

- ↳ CPU execute multiple jobs
- ↳ switches so quickly that user interaction with programs is possible
- ↳ allows many users to share the comp simultaneously
- ↳ CPU scheduling
- ↳ processes are in CPU
- ↳ multiple processes running concurrently and preempted after a fixed time slice



- ↳ increases CPU utilization
- ↳ doesn't leave CPU idle

OPERATING SYSTEM SERVICES

CHP 2

One set of operating-system services provides functions that are helpful to the user:

- **User interface** - Almost all operating systems have a user interface (UI).
- Varies between Command-Line (CLI), Graphics User Interface (GUI), Batch

○ **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)

○ **I/O operations** - A running program may require I/O, which may involve a file or an I/O device

○ **File-system manipulation** - The file system is of particular interest. Programs need to read and write files and directories, create and delete them, search them, list file information, permission management.

○ **Communications** - Processes may exchange information, on the same computer or between computers over a network

- Communications may be via shared memory or through message passing (packets moved by the OS).

▪ **Error detection** - OS needs to be constantly aware of possible errors

- May occur in the CPU and memory hardware, in I/O devices, in user programs
- For each type of error, OS should take the appropriate action

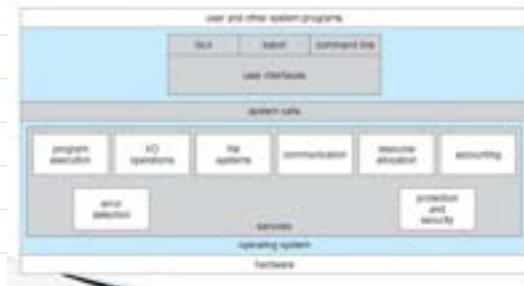
▪ **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them

- Many types of resources - CPU cycles, main memory, file storage, I/O devices.

▪ **Accounting** - To keep track of which users use how much and what kinds of computer resources

▪ **Protection and security** - The owners of information stored in a multiterminal or networked computer system may want to control use of that information; concurrent processes should not interfere with each other

- **Protection** involves ensuring that all access to system resources is controlled.
- **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts



CLI

↳ type

↳ interpreters are known as shells

IF PEARS PVC

ELPC

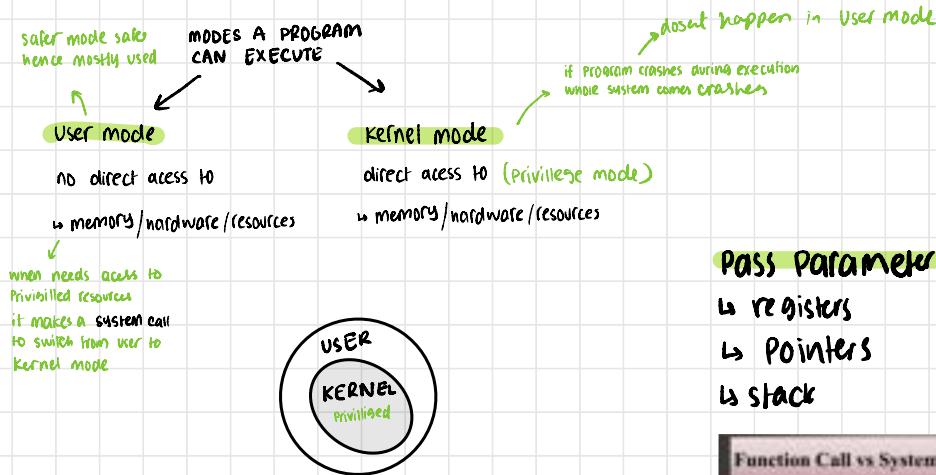
GUI

↳ icons

SYSTEM CALLS

System calls

- ↳ Provide an interface to the services made available by an OS in kernel mode
 - ↳ Made by programs to access certain resources



TYPES OF SYSTEM CALLS

1 Process Control

- execute
 - create process
 - fork on Unix-like systems or
 - NoCreateProcess in the Windows
 - terminate process
 - get/set process attributes
 - wait for time, wait event, signal event
 - allocate free memory

2. File management

- create file, delete file
 - open, close
 - read, write, reposition
 - get/set file attributes

Device Management

 - request device, release device
 - read, write, reposition
 - get/set device attributes
 - logically attach or detach devices

3. Device Management

- request device, release device
 - read, write, reposition
 - get/set device attributes
 - logically attach or detach devices

Information Maintenance

- get/set time or date
 - get/set system data
 - get/set process, file, or device attributes

5 Communication

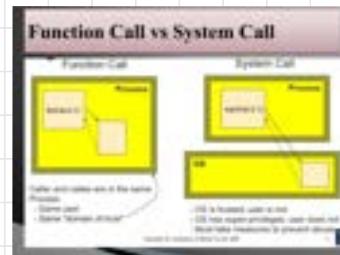
 - create, delete communication connection
 - send, receive messages
 - transfer status information
 - attach or detach remote device

Common mistakes

- create, delete communication connection
 - send, receive messages
 - transfer status information
 - attach or detach remote device

6. Security

- Protection
 - Control access to resources
 - Get and set permissions
 - Allow and deny user access



DIPS FC

FIPC

SYSTEM PROGRAMS

 **File management** - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories

 **Status information**

- Some ask the system for info - date, time, amount of available memory, disk space, number of users
- Others provide detailed performance, logging, and debugging information

Background Services

- Launch at boot time
- Provide facilities like disk checking, process scheduling, error logging, printing
- Run in user context not kernel context
- Known as *services, subsystems, daemons*

Application programs

- Run by users

File modification

- Text editors to create and modify files
- Special commands to search contents of files or perform transformations of the text

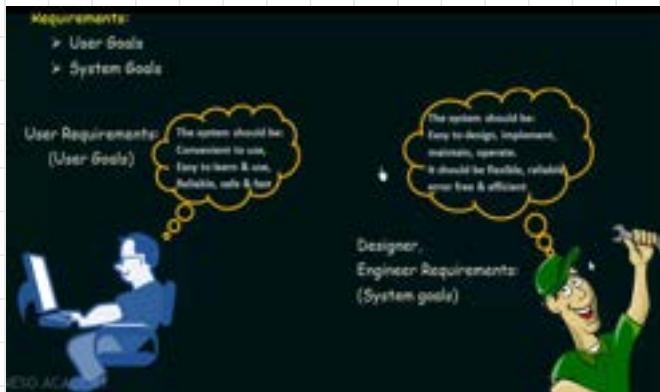
 **Programming-language support** - Compilers, assemblers, debuggers and interpreters sometimes provided

 **Program loading and execution** - debugging systems for higher-level and machine language

 **Communications** - Provide the mechanism for creating virtual connections among processes, users, and computer systems

FIPC

OS DESIGN and IMPLEMENTATION



Mechanisms and Policies

Mechanisms determine how to do something

Policies determine what will be done

One important principle is the separation of policy from mechanisms.

OS STRUCTURES

1) SIMPLE STRUCTURE

- ↳ if one program fails entire system crashes
- ↳ not well protected / structure



2) MONOLITHIC STRUCTURE

- ↳ too many functions packed in one level → kernel part
- ↳ hence implementation/maintenance difficult
- ↳ to change one thing will have to touch entire kernel



3) LAYERED STRUCTURE

pros

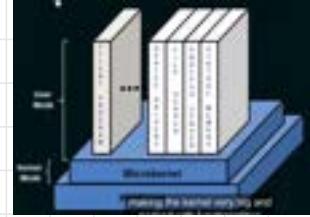
- ↳ functionalities not packed into same layer
- ↳ easy to implement/debug
- ↳ if one layer problem then only have to look at that layer
- ↳ hardware protected by layers above / no direct access by the user
- ↳ a layer can only use those below the layer
- ↳ not efficient
- ↳ when a layer wants to use services of below layer it will have to go down one by one
- ↳ may not be fast



4) MICROKERNELS STRUCTURE

pros

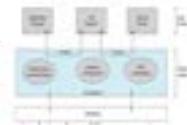
- ↳ micro kernel provides core functionalities
- ↳ other services are implemented as a user level program
- ↳ communication b/w client programs and system programs as made through message passing
- ↳ can continue running mostly in user mode so less chances of system crash
- ↳ due to message passing



→ best

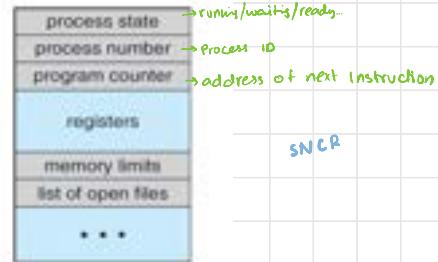
5) MODULE STRUCTURE

- ↳ uses OOP techniques
- ↳ core kernel has core functionalities
- ↳ other functionalities are in the form of modules
- ↳ modules loaded into kernel when required
- ↳ doesn't have to go through all layers like in layer approach
- ↳ each module can communicate with other module directly
- ↳ no need for message passing hence no performance decrease



CHP 3

PCB

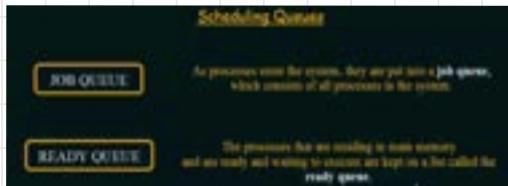


SNCR

Process Scheduling

- > The objective of multiprogramming is to have some process running at all times, to maximize CPU utilization.
- > The objective of time sharing is to switch the CPU among processes as frequently that users can interact with each program while it is running.
- > To meet these objectives, the process scheduler selects an available process (possibly from a set of several available processes) for program execution on the CPU.

Process alternate b/w these 2 states



- I/O-bound process - spends more time doing I/O than computations, many short CPU bursts
- CPU-bound process - spends more time doing computations; few very long CPU bursts

CONTEXT SWITCH

↳ current program

↳ P_1 executing

↳ interrupt occurs

Present in PCB
of the process

↳ system saves state of P_1 context

↳ CPU switches to P_2 and suspends P_1

↳ P_2 executes and finishes

↳ restore state of P_1

↳ resume execution of P_1

Pure overhead → resources/time spent...
cost that is involved in doing something

↳ does no useful work while switching

Kernels actions in context switch

↳ Saving the current process state

↳ selecting the next process

↳ loading the next process state

↳ restoring saved process

↳ handing control to next process

↳ resuming execution

fork():

↳ creates a separate duplicate process

↳ Parent ^{its duplicate} → child

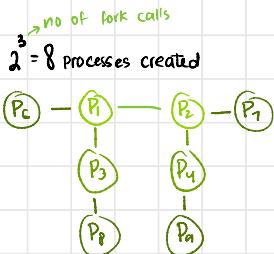
↳ ID is different

exec():

↳ program in parameter will replace entire process

↳ ID remains same

main() {
 fork()
 fork()
 fork()
 }
 ↳ no of fork calls
 $2^3 = 8$ processes created



Interrupt

cause OS to change CPU from current task and to run a kernel routine

CHPS

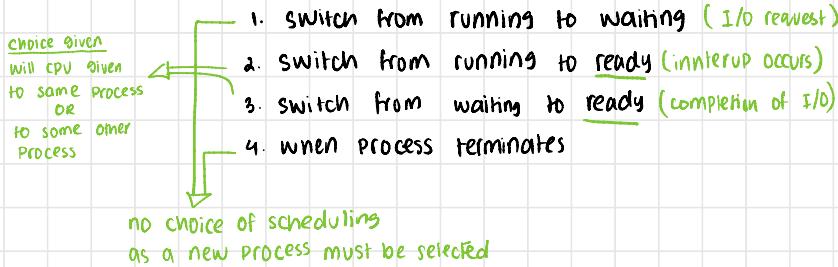
CPU SCHEDULER

- ↳ whenever CPU becomes idle
CPU scheduling
- ↳ OS selects a process in ready queue to execute

Dispatcher

- ↳ giving control to process selected by short term scheduler (CPU scheduler)
- ↳ time to stop one process and start another is dispatch latency

CPU scheduling circumstances



2 WAYS CPU SCHEDULING TAKES PLACE

Preemptive scheduling

- ↳ circumstance 2 and 3
- ↳ CPU can be taken away from a process before it has went into ready state

non Preemptive scheduling/cooperative

- ↳ circumstance 1 and 4
- ↳ CPU will never be taken away from a process until and unless it has completed execution (terminated)
- ↳ went into waiting state

CONS

shared memory

- ↳ P₁ writing a sentence in shared memory
- ↳ P₂ came, P₁ preempted/halted
- ↳ CPU given to P₂
- ↳ P₂ when reads from shared memory
- ↳ P₂ reads inconsistent data as P₁ wasn't done writing

FCFS (First Come First Serve)

- when process enters ready queue
- PCB linked to tail of the queue
- avg waiting time long
- non preemptive

turnaround time = Completion time - Arrival time

Waiting time = turnaround time - Burst time

CONS

- disasterous to allow one process CPU for extended time
- CONVOY EFFECT
 - higher burst time arrive before smaller
 - smaller process have to wait long time

Process	Burst Time ms
P ₁	24
P ₂	3
P ₃	3

Process ID	Arrival Time	Burst Time
P ₁	0	24
P ₂	6	3
P ₃	12	3
P ₄	18	2
P ₅	24	4



Process ID	Completion Time	Turnaround Time	Waiting Time
P ₁	8	8 - 0 = 8	0 - 0 = 0
P ₂	17	17 - 6 = 11	6 - 6 = 0
P ₃	17	17 - 12 = 5	12 - 12 = 0
P ₄	19	19 - 18 = 1	18 - 18 = 0
P ₅	23	23 - 24 = 8	24 - 24 = 0

$$\text{avg turnaround time} = \frac{8+11+5+1+8}{5} = 8 \text{ units}$$

$$\text{Waiting time} = 0 + 24 + 27 = 51 \text{ ms}$$

$$\text{avg waiting time} = \frac{51}{3} = 17 \text{ ms}$$

$$\text{avg waiting time} = \frac{0+7+0+11+4}{5} = 4.4 \text{ units}$$

SJF (shortest job first)

- ↳ If CPU burst are same FCFS is used to break tie
- ↳ Preemptive or non preemptive

CONS

- ↳ No way to know length of next CPU burst

turnaround time = completion time - arrival time

Waiting time = total waiting time - ms process executed - Arrival time

SOLUTION

- ↳ We can approximate
- ↳ We can predict using previous CPU bursts

NON PREEMPTIVE

Process	Burst Time
P ₁	24
P ₂	3
P ₃	3

PREEMPTIVE

Process ID	Arrival Time	Burst Time
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5



→ as no arrival time given assume 0

$$\text{Waiting time} = 0 + 3 + 6 = 9 \text{ ms}$$

$$\text{avg waiting time} = 9/3 = 3 \text{ ms}$$

$$\text{P}_1 \quad \text{waittime} = (10-1-0) + (1-0-1) + (17-0-2) + (5-0-3)$$

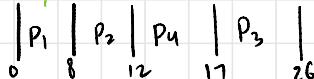
$$\text{avg} = 26/4 = 6.5$$

$$\text{turnaround time} = (10-0) + (5-1) + (26-2) + (10-3)$$

$$\text{avg} = 52/4 = 13$$

NON PREEMPTIVE

→ can't end in mid as preemph



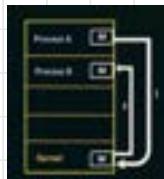
INTER PROCESS COMMUNICATION (IPC)

- ↳ Mechanism for processes executing concurrently to communicate and synchronise their actions

MESSAGE PASSING

↳ TWO OPERATIONS

- ↳ send message → fixed size message/variable
- ↳ receive message



IN ORDER TO COMMUNICATE USING MESSAGE PASSING

- ↳ establish a communication link b/w them
 - ↳ Physical (shared memory/hardware bus)
 - ↳ logical (logical properties)
- ↳ exchange messages via send/receive

Priority Scheduling

- ↳ CPU allocated by priority

CON: Starvation → low priority may never execute

SOLUTION: Aging → increase priority as time progresses

- ↳ turn around time = completion - arrival time
- ↳ Waiting time = Total waiting time - ms process executed - Arrival time
Preemptive
- ↳ Waiting time = turn around time - burst time
non preemptive

PREEMPTIVE

Process ID	Arrival Time	Burst Time	Priority
P1	0	/a	2
P2	5	20	0
P3	12	2	1
P4	2	/ 7	3
P5	9	16	4

Arrival	Completion	Waiting Time
0	0	0
5	20	15
12	22	10
2	49	47
9	65	56

NON PREEMPTIVE

$$\text{Waiting time} = (40-2-0) + (5-0-5) + (49-0-12) + (33-3-2) + (51-0-9)$$

$$\text{avg} = 29/5$$

→ best for time sharing

Round Robin (RR) Scheduling

↳ small time slices for each process (time quantum)

↳ FCFS but has preemption



2 Possibilities

↳ burst time < time quantum

↳ burst time > time quantum

↳ CPU released

↳ timer goes off

↳ CPU scheduler gets next process in ready queue

↳ interrupt to OS
↳ context switch, process moved to tail of ready queue

↳ CPU scheduler gets next process in ready queue

FIFO in ready queue

Tail → new processes

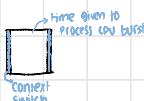
Head → next process to get CPU

CONS

time quantum

↳ too small: too many context switches

↳ too big: starve for long time



Process ID	Burst Time
P1	24
P2	3
P3	5



Method 1:

Turn Around Time = Completion Time - Arrival Time

Waiting Time = Turn Around Time - Burst Time

Process	Completion Time	Turn Around Time	Waiting Time
P1	36	36 - 0 = 36	36 - 24 = 12
P2	7	7 - 0 = 7	7 - 3 = 4
P3	18	18 - 0 = 18	18 - 5 = 13

Average Turn Around Time
 $= (12 + 4 + 13) / 3$
 $= 47 / 3 = 15.67 \text{ ms}$

Average waiting time
 $= (12 + 4 + 7) / 3$
 $= 23 / 3 = 7.67 \text{ ms}$

Method 2:

Waiting Time = Last Start Time - Arrival Time - (Preemption * Time Quantum)

Process	Waiting Time
P1	$36 - 0 - (24 \times 1) = 0$
P2	$4 - 0 - (24 \times 1) = 4$
P3	$7 - 0 - (24 \times 1) = 7$

Average waiting time
 $= (0 + 4 + 7) / 3$
 $= 11 / 3 = 3.67 \text{ ms}$

ishma hafeez notes

represent

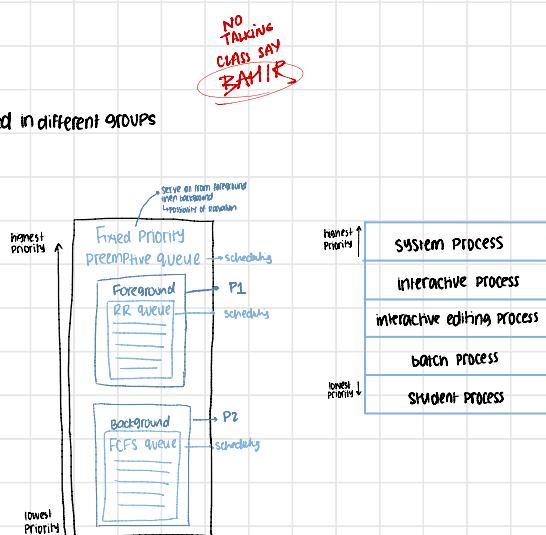
Multi-level Queue Scheduling

Scheduling algorithms for situations in which processes are classified in different groups

Eg

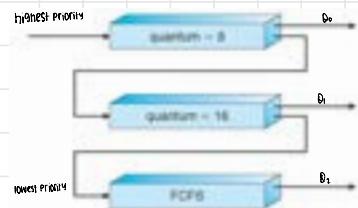


- ↳ Partitions ready queue into several separate queues
- ↳ Scheduling takes place within and among these queues
- ↳ Each process permanently assigned to 1 queue
based on memory size, priority and process type
- ↳ Each queue has its own scheduling algorithm (FCFS, SJF, RR...)



Multi-level Feedback Queue Scheduling

- ↳ Allows processes to move between queues
- ↳ Separate processes according to their CPU bursts
- ↳ If a process uses too much CPU time
then moved to lower priority queue (so everybody gets a chance)
- ↳ I/O and interactive processes are in higher priority queues
as they need quick responses
- ↳ A process in lower priority queue for too long → aging
may be moved to higher priority queue (Prevents starvation)



Scheduling

- A new job enters queue Q₀, which is served FCFS
 - * When it gains CPU, job receives 8 milliseconds
 - * If it does not finish at 8 milliseconds, job is moved to queue Q₁
- At Q₁, job is again served FCFS and receives 16 additional milliseconds
 - * If it still does not complete, it is preempted and moved to queue Q₂

Parameters

- ↳ No. of queues
- ↳ Scheduling algo for each queue
- ↳ Method to upgrade process to higher priority
- ↳ Method to demote process to lower priority
- ↳ Method to determine queue when process needs service
demote/upgrade

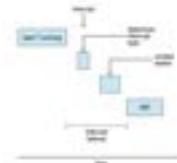
MULTI-PROCESSOR SCHEDULING

- CPU scheduling more complex when multiple CPUs are available
- **Homogeneous processors** within a multiprocessor
- **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
- **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
- **Processor affinity** – process has affinity for processor on which it is currently running
- **soft affinity**: When an operating system has a policy of attempting to keep a process running on the same processor—but not guaranteeing that it will do so known as **soft affinity**.
- **hard affinity**: allowing a process to specify a subset of processors on which it may run.

NOTES
MISSING

Real-Time CPU Scheduling

- **Hard real-time systems** – no guarantees as to when individual real-time process will be scheduled
- **Soft real-time systems** – real-time is guaranteed by deadline
- **Time-triggered architecture** (periodic tasks)
- **Interrupt latency** – time from arrival of interrupt to start of routine that service interrupt
- **Dispatch latency** – time for scheduler to take control process off CPU and switch to another



Multiple-Processor Scheduling – Load Balancing

- If SMP, need to keep all CPUs loaded for efficiency
- **Load balancing** attempts to keep workload evenly distributed
- **Push migration** – periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs
- **Pull migration** – idle processors pulls waiting task from busy processor

CONFUSED

PROCESS CONCEPT

- ↳ Process are executed programs that have

↳ Resource Ownership unit → process/task
 ↳ process includes virtual space? to hold process image

- ↳ OS prevents unwanted interference b/w processes

Scheduling / Execution

- Process follows an execution path that may be interleaved with other processes
- Process has an execution state (Running, Ready, etc.) and a scheduling priority and is scheduled and dispatched by the operating system
- Today, the unit of dispatching is referred to as a thread or lightweight process

MOTIVATION

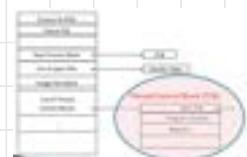
- ↳ Most modern applications are multithreaded
- ↳ Threads run within application (process)
- ↳ Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- ↳ Process creation is heavy-weight while thread creation is light-weight
- ↳ Can simplify code, increase efficiency
- ↳ Kernels are generally multithreaded

Thread control block

- ↳ Information associated with each thread
- ↳ Program Counter
- ↳ CPU Registers
- ↳ CPU Scheduling info
- ↳ Pending I/O Info

Process control block (PCB)

- ↳ Information associated with each process
- ↳ Memory management info
- ↳ Accounting info



Thread operations

- ↳ Spawn: a thread within a process may spawn another thread
 - ↳ Provides new thread instruction pointer, arguments, register, stack
- ↳ Block: a thread needs to wait for an event
 - ↳ saves its user registers, program counter, stack pointers
- ↳ Unblock: when the event for which the block occurs
- ↳ Finish: when thread completes
 - ↳ its register context and stacks are deallocated

Thread states

- ↳ running
- ↳ ready
- ↳ blocked

Multithreaded Applications

- ↳ An application that creates photo-thumbniling from a collection of images may use a separate thread to generate a thumbnail from each image image.
- ↳ A multi-threaded application may have one thread display images or two while another thread retrieves data from the network.
- ↳ A word processor may have a thread for displaying graphics, another thread for responding to keystrokes from the user, and a third thread for performing spelling and grammar checking in the background.

LINUX Kernel is also multithreaded

Single vs Multithreaded Webserver

- ↳ How a web-server (as a single process) increase its response time?

One solution is to have the server run as a single process that accepts requests. When the server receives a request, it creates a separate process to service that request. In fact, this process-creation method was in common use before threads became popular. Process creation is time consuming and resource intensive, however. If the new process will perform the same tasks as the existing process, why incur all that overhead?

Single vs Multithreaded Webserver

- ↳ Answer: 2009 - request per second.
- ↳ Performance of single threaded server
- ↳ Working of a multi-threaded server
- ↳ What benefits do we get using a multithreaded server?



Process

- ↳ A program in execution
- ↳ can have 1 or more threads



Threads

- ↳ unit of execution within a process

- ↳ It compromises of

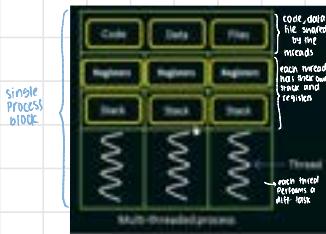
1. Thread ID
2. Program Counter
3. Register set
4. Stack

↳ shares code section, data section, OS resources
with other threads of same processes

	Process	Thread
1.	Process is basic unit of resource allocation.	Thread is light weight unit of resource allocation than a process.
2.	Process controlling access interaction with operating system.	Thread switching does not need to interact with operating system.
3.	In multiprocess environment, each process executes the same code but has its own memory and file resources.	All threads can share same set of physical memory, called processes.
4.	If one process is blocked then no other process can execute until that thread is unblocked.	One thread is blocked while waiting, another thread in other core can run.
5.	Multiprocess environment using threads can share resources.	Advantage: threads can share common resources.
6.	No multiprocess environment requires separate copy of memory of file system.	One thread can read, write or change another thread's state.

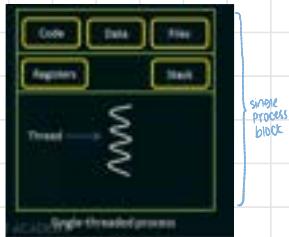
Multi Thread

- ↳ performs multiple tasks at a time



Single Thread

- ↳ performs one task at a time



MULTI-THREADING BENEFITS

- ↳ **Responsiveness:** may continue execution if process is blocked / performing lengthy operation
- ↳ **Dedicated threads for handling user events**
- ↳ **Resource Sharing:** threads share memory and resources of a process → allows diff threads activity within same address space
 - ↳ better than shared memory / message passing
 - ↳ avoiding interprocess contention
- ↳ **Economy:** cheaper than process creation as threads share resources
- ↳ **Scalability:** utilization of multiple cores for parallel execution
 - ↳ increases concurrency

Multicore Programming



parallel execution

PARALLELISM

↳ act of managing multiple computations simultaneously



DATA PARALLELISM ← IMP → TASK PARALISM

↳ focus on distributing data

across diff parallel computing nodes

↳ focus on distributing threads

across diff parallel computing nodes

Data Parallelism
Same operations are performed on different subsets of same data.

Synchronous computation

Speedup is more as there is only one execution thread operating on all sets of data.

Amount of parallelization is proportional to the input data size.

Designed for optimum load balance on multi processor system.

Task Parallelism
Different operations are performed on the same or different data.

Asynchronous computation

Speedup is less as each processor will execute a different thread or process on the same or different set of data.

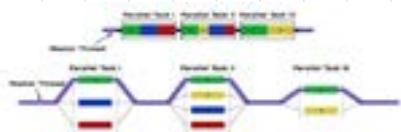
Amount of parallelization is proportional to the number of independent tasks to be performed.
Load balancing depends on the availability of the hardware and scheduling algorithms like static and dynamic scheduling.

CONCURRENCY

↳ act of managing multiple computations at the same time

but not simultaneously → control is switched

FORK-JOIN MODEL



advantages:
If problem is small enough:
solve problem directly (sequential algorithm)
else:
For part in subproblem:
fork subtask to solve part
join all subtasks opened to previous step
combine results from subtasks

↳ Multicore systems putting pressure on programmers, challenges include:

Dividing activities

- What tasks can be separated to run on different processors

Balance

- Balance work on all processors

Data splitting

- Separate data to run with the tasks

Data dependency

- Watch for dependences between tasks

Testing and debugging

- Harder!!!!

AMDAHL'S LAW

↳ Speed up in latency of a task execution

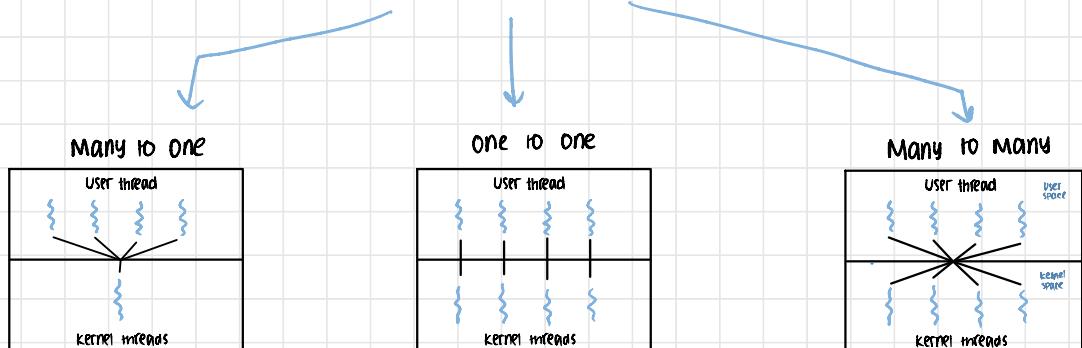
$$\text{Speed UP} \leq \frac{1}{S + \frac{S(1-S)}{N \cdot \text{no of cores}}}$$

B) $S=25\%, N=2$

$$\text{Speed UP} \leq \frac{1}{0.25 + \frac{0.25}{2}} = 2.28$$

MultiThreading → multiple threads at the same time

establish relationship between user and kernel thread



PROS

- ↳ multiple user threads mapped to 1 kernel thread → efficient
- ↳ thread management is done by threaded library

CONS

- ↳ if 1 thread makes a blocking system call → entire process will be blocked
- ↳ only 1 thread can access kernel at a time.
- ↳ multiple threads are unable to run in parallel on multiprocessors

PROS

- ↳ maps each user thread to kernel thread ↳ always another thread to run when a thread makes a blocking system call
- ↳ more concurrency
- ↳ allows multiple threads to run parallel on multi processors

CONS

- ↳ each user thread requires a corresponding kernel thread
- ↳ overhead of creating kernel threads burdens performance
- ↳ restricts no. of threads supported by the system
 - ↳ e.g. 4 core processor
 - ↳ 5 threads
 - ↳ then only 4 threads work parallel
 - ↳ hence no. of threads restricted

PROS

- ↳ kernel threads are ≤ to user threads
- ↳ more concurrency
- ↳ no limit to creating user threads
- ↳ kernel threads can run parallel on a multi processor

Hyper Threading / SMT (Simultaneous multithreading)

- ↳ more than 1 multithreading going on in the same system

e.g. ↳ 4 core → virtually/ logically divided into multiple processors
↳ 8 threads support

TYPES OF THREADS

User threads

- ↳ implemented by user
- ↳ requires no hardware support
- ↳ easy to implement
- ↳ if one thread blocked, entire process blocked

Kernel threads

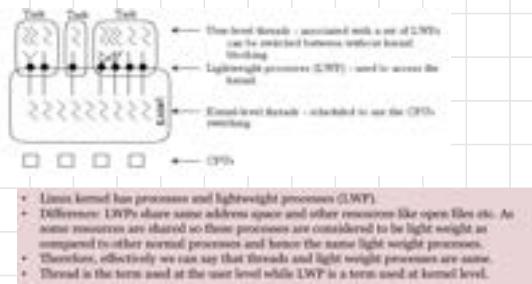
- ↳ implemented by kernel
- ↳ requires hardware support
- ↳ difficult to implement
- ↳ if one thread blocked, another thread continues execution

POSIX PThreads

- ↳ Win32 Threads
- ↳ Java Threads

POSIX Pthreads

- May be provided either as user-level or kernel-level
 - A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
 - API specifies behavior of the thread library, implementation is up to development of the library
- Win32**
- Kernel-level library on Windows system
 - Java
 - Java threads are managed by the JVM
 - Typically implemented using the threads model provided by underlying OS



Threaded Libraries

$$\text{SUM} = \sum_{i=1}^N i$$

APPLICATION
LIBRARY
KERNEL

→ user function() → wrapped and threaded

Asynchronous Threading

- ↳ Once parent creates child thread, parent resumes execution
- ↳ They execute concurrently and independently of one another
- ↳ Threads are independent → no dependency
- ↳ little data sharing

Synchronous Threading

- ↳ Once parent creates 1 or more children
- ↳ it waits all child to terminate before it resumes
- ↳ threads by parent work concurrently
- ↳ significant data sharing among threads

Race condition

- ↳ when several processes manipulate / access the same data concurrently
- and outcome depends on order in which access take place
- results in inconsistency

(SOLUTION)

Process Synchronization

The orderly execution of cooperating process that share an address space

Producer and consumer run concurrently using
BUFFER

Unbounded buffer

- ↳ has a size on buffer
- ↳ consumer waits for new items
- ↳ producer can always produce new items

bounded buffer

- ↳ fixed buffer size
- ↳ consumer must wait if buffer empty
- ↳ producer must wait till buffer full
- ↳ use counter variable
 - + produced
 - consumed

Critical section problem

- ↳ a system with n processes
- ↳ each process has a segment of code called critical section

↳ where process may be changing common variables
updating table
writing a file

Three requirements

- ↳ mutual exclusion:

If process P_i is executing in its critical section, then no other processes can be executing in their critical sections.

- ↳ progress:

If no process is executing in its critical section and some processes wish to enter their critical sections, then only those processes that are not executing in their remaining sections can participate in the decision on which will enter its critical section next, and this selection cannot be postponed indefinitely.

- ↳ bounded waiting:

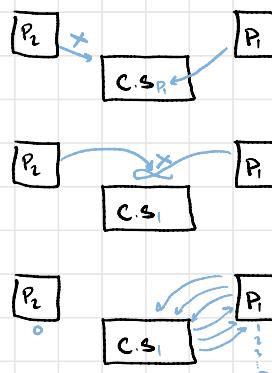
There exists a bound, or limit, on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.

Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- All solutions follow based on idea of locking:
 - ↳ Protecting critical sections via locks
- Implementations could disable interrupts:
 - ↳ Currently running code would execute without preemption.
 - ↳ Generally too inefficient on multiprocessor systems.
 - ↳ Operating systems using them not broadly suitable
- Modern machines provide special atomic hardware instructions:
 - ↳ Atomic → non interruptible
 - ↳ Either load memory word and set value
 - ↳ Or swap contents of two memory words

Two approaches depending on if kernel is preemptive or non-preemptive.

- Preemptive → allows preemption of process when running in kernel mode
- Non-preemptive → runs until exits kernel mode, blocks, or voluntarily yields CPU
 - ↳ Essentially free of race conditions in kernel mode.



Peterson's SOLUTION

→ humble

↳ software solution to the critical section problem

↳ 2 processes that alternate execution

↳ b/w
critical section,
remainder section

↳ The 2 processes share 2 variables

↳ int turn;

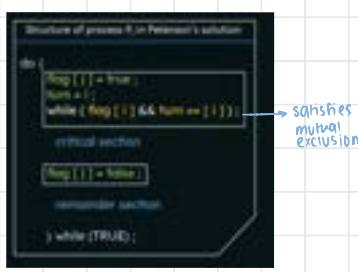
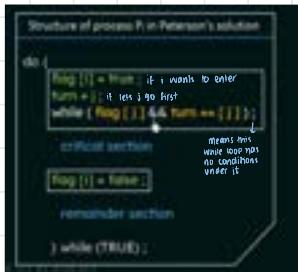
↳ indicates turns who
is to enter critical section

↳ boolean flag[2];

↳ indicates if process

is ready to enter critical section

↳ Satisfies all 3 critical section requirements



Test and Set Lock

↳ hardware solution to the critical section problem

↳ Processes share shared lock variable = 0 → unlocked
1 → locked

↳ helps satisfy
mutual exclusion

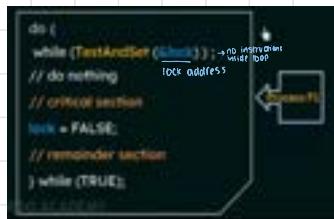
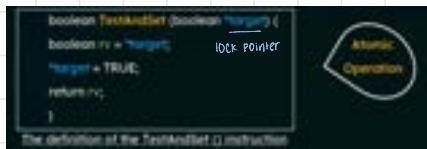
→ lock=0 in beginning

↳ If lock=1, then wait till it becomes free

↳ If lock=0, then execute critical section and lock=1

↳ Satisfies mutual exclusion

↳ does not satisfy bounded-waiting



Atomic operation

↳ single, uninterrupted operation

e.g.

TestAndSet() is run as a single operation
which can not be interrupted

Bounded-waiting Mutual Exclusion with test_and_set

```

do {
    mutexing[i] = name;
    key = true;
    while (mutexing[i] == key) {
        key = !mutexing[i];
        mutexing[i] = false;
        /* critical section */
        g = i + j - k - m;
        while (g >= i) {
            g = g - i;
            k = k - i;
            lock = false;
        }
        lock = false;
    }
    /* remaining sections */
} while (true);

```

Compare and Swap Instruction

```

int compare_and_swap(int *value, int expected, int new_value) {
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}

while (true) {
    while (compare_and_swap(&v, 0, 1) == 0)
        /* do nothing */
    /* critical section */
    lock = 1;
    /* remainder section */
}

```

If two CAS instructions are executed simultaneously (each on a different core), they will be executed sequentially in some arbitrary order.

On Intel x86 architectures
lock acquire: **lock** **sfence**
lock release: **sfence**

MUTEX LOCKS/SPIN LOCK

↳ to protect critical section/prevent race conditions

↳ **acquire()** } both calls must
↳ **release()** } be atomic

↳ requires busy waiting

while PROCESS in critical section
any other process that tries to enter critical section
must loop continuously to check if condition true

↳ wastes CPU cycles

↳ that some other process could use it concurrently

- **acquired()** :


```

while (!available)
    /* Busy wait */
    available = false;
      
```
- **released()** :


```

available = true;
      
```
- **do {**

```

acquire lock
critical section
release lock
remainder section
      
```
- **while (true);**

Pthreads Synchronization – Mutex Lock

```

#include <pthread.h>
pthread_mutex_t mutex;

/* create the mutex lock */
pthread_mutex_init(&mutex, NULL);

1st Arg: Mutex instance
2nd Arg: Attributes, NULL means no error checks will be performed

```

Pthreads Synchronization – Mutex Lock

- The mutex is acquired and released with the **pthread_mutex_lock()** and **pthread_mutex_unlock()** functions.
- If the mutex lock is unavailable when **pthread_mutex_lock()** is invoked, the calling thread is blocked until the owner invokes **pthread_mutex_unlock()**.

```

pthread_t tid[2];
int counter;
pthread_mutex_t lock;

```

```

main() {
    /* Create threads */
    pthread_create(&tid[0], NULL, &counter, &lock);
    pthread_create(&tid[1], NULL, &counter, &lock);
    /* Wait for threads to start */
    sleep(1);
    /* Set lock to be unavailable */
    pthread_mutex_lock(&lock);
    /* Release lock */
    pthread_mutex_unlock(&lock);
}

```

```

for (i = 0; i < 10; i++) {
    /* Create threads */
    pthread_create(&tid[0], NULL, &counter, &lock);
    pthread_create(&tid[1], NULL, &counter, &lock);
    /* Wait for threads to start */
    sleep(1);
    /* Set lock to be unavailable */
    pthread_mutex_lock(&lock);
    /* Release lock */
    pthread_mutex_unlock(&lock);
}

```

<http://www.pthreadguide.org/pthreads-lock-for-mutual-thread-synchronization>

SEMAPHORE

↳ software solution to the critical section problem

↳ technique to manage concurrent processes

↳ an int variable S shared b/w threads
↳ always true
↳ semaphore

↳ S can only be accessed through 2 atomic operations

↳ $\text{wait}()$ → to decrement

$S \leq 0$ → some process already in critical section → wait
 $S > 0$ → can use shared resource → don't wait

↳ $\text{signal}()$ → to increment

tell other processes that shared resource is free to be used

```
Definition of wait(L)
P(Semaphore S) {
    while (S == 0)
        ; // no operation
    S--;
    /* when S>0 */
}
```

```
Definition of signal(L)
V(Semaphore S) {
    S++;
}
```

* $S = 1$ in beginning

TWO TYPES OF SEMAPHORES

Binary Semaphore

↳ S can only have 2 values 0 or 1

↳ Some process already executing
↳ not free to use shared resource
↳ provide mutual exclusion

↳ behave similarly to mutex locks

↳ can implement a counting semaphore S as a binary semaphore

Counting Semaphore

↳ S can multiple values

↳ used to control access to a resource that has multiple instances

↳ $S = \text{no. of instances of a resource}$

SEMAPHORE DISADVANTAGE

↳ requires BUSY WAITING

↳ while a process is in critical section
any other process that tries to enter critical section
must loop continuously to check if condition true

↳ wastes CPU cycles

↳ that some other process could use
productively

* This type of semaphore also known as SPINLOCK

PROBLEM

↳ Deadlock

↳ set of blocked processes, each holding a resource
and waiting to acquire a resource, held by another process

↳ Starvation

↳ when a process is postponed because it requires a resource
to run, that is never allocated to this process

SOLUTION TO BUSY WAITING

↳ modify $\text{wait}()$ and $\text{signal}()$

↳ waiting state

↳ to block() → places process in waiting queue

↳ to wake up() → move process from
waiting to ready queue

↳ instead of busy waiting, block itself

↳ transfers control to CPU scheduler
↳ which selects another process
to execute

↳ hence no wastage of CPU time



DEADLOCK

↳ process stuck in circular waiting for the resources

VS

STARVATION

↳ process waits for a resource indefinitely

→ deadlock implies starvation BUT starvation does not imply deadlock



POSIX - Semaphores

```
#include <semaphore.h>
Sem_t sem;
/* Create the semaphore and initialize it to 1 */
Sem_init(&sem, 0, 1);
The sem_init function is passed three parameters:
1. A pointer to the semaphore
2. A flag indicating the level of sharing
3. The semaphore's initial value
```

POSIX - Semaphores

```
/* acquire the semaphore */
Sem_wait(&sem);
/* critical section */
/* release the semaphore */
Sem_post(&sem);
```

Priority Inversion

↳ scheduling problem when

lower Priority process holds a lock

needed by a higher Priority process

SOLUTION



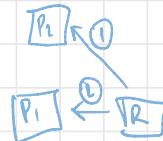
Priority Inheritance Protocol

↳ when a job blocks a higher priority job

it ignores the current lower priority job

executes higher priority job's critical section

then releases lock and returns to the current lower priority job



CLASSICAL PROBLEMS OF SYNCHRONIZATION

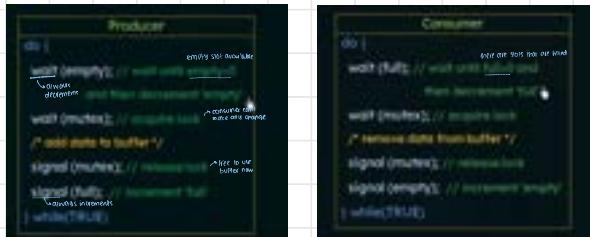
1. Bounded buffer Problem

- ↳ shared buffer
 - ↳ buffer of n slots, each capable of storing one unit of data
 - ↳ producer process
 - ↳ insert data in empty slot of buffer
 - ↳ not insert data when buffer is full
 - ↳ consumer process
 - ↳ remove data from filled slot of buffer
 - ↳ not remove data when buffer is empty



Solution using semaphores

- ↳ n buffers \rightarrow each can hold one item
 - ↳ $m(mutex)$ \rightarrow binary semaphore
 - ↑ acquire lock
 - ↓ release lock
 - ↳ empty \rightarrow counting semaphore
 - ↑ initial value = no of slots in buffer
 - ↓ since nothing in buf, empty
 - ↳ full \rightarrow counting semaphore
 - ↑ initial value =

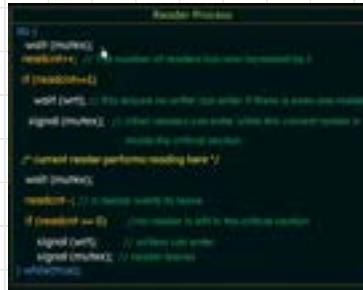


2. The Readers writers Problem

- ↳ shared database
 - ↳ Readers → only want to read
 - ↳ Writers → want to read and write
 - ↳ if a writer and a writer/reader access database simultaneously → PROBLEM
 - ↳ give exclusive access to database → SOLUTION

Solution using semaphores

- ↳ **mutex** → a semaphore → initialized to 1
 - ↳ **wrt** → a semaphore → initialized to 1
 - ↳ **readcount** → counts how many readers reading same data
 - ↳ whenever modified wait(mutex) to be called



→ used in OS resource allocation → Processes
→ resources

3. The Dining-Philosophers Problem

↳ 5 philosophers, 5 forks

↳ Eating
use 2 forks for eating
pick 1 fork on a time
can eat unless
has 2 forks

↳ Thinking → idle

↳ When a philosopher eats he uses 2 forks

↳ NO 2 adjacent philosophers try to eat at the same time → SOLUTION

SOLUTION USING SEMAPHORES

↳ fork[5] → a binary semaphore array
as 5 forks
all element initialized to 1
→ i-free
→ 0 = free

↳ Grab fork → wait() → used on 2 forks

↳ release fork → signal()



can lead to deadlock

All 5 hungry and grab left fork

all elements of fork = 0

when grabbing right fork, they will all be delayed forever

SOLUTIONS TO DEADLOCK

→ no of forks will remain 5

1. ALLOW AT MOST 4 PHILOSOPHERS

2. ALLOW TO PICK FORKS ONLY IF BOTH FORKS AVAILABLE

3. USE ASYMMETRIC SOLUTION

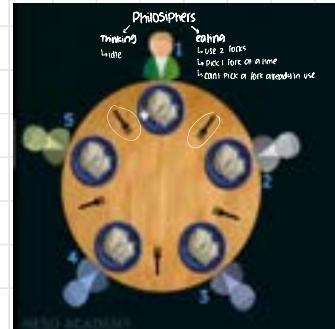
↳ an odd P first picks right fork then left fork

↳ an even P first picks left fork then right fork

→ deadlock implies starvation

B U T

starvation does not imply deadlock



ishma hafeez
notes

represent

DEADLOCKS

- ↳ SET OF BLOCKED PROCESSES, EACH HOLDING A RESOURCE AND WAITING TO ACQUIRE A RESOURCE, HELD BY ANOTHER PROCESS

IF ANY 1 NOT THERE
NO DEADLOCK SITUATION

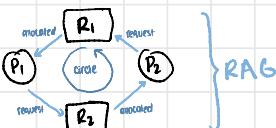
4 Deadlock necessary conditions

1. Mutual Exclusion no resource can be used by more than 1 process at a time

2. NO Preemption 1 process is holding a resource and second process comes and try to preempt resource to 1st, this shouldn't occur

3. Hold and Wait P1 holding R1, wants R2;
P2 holding R2, wants R1.

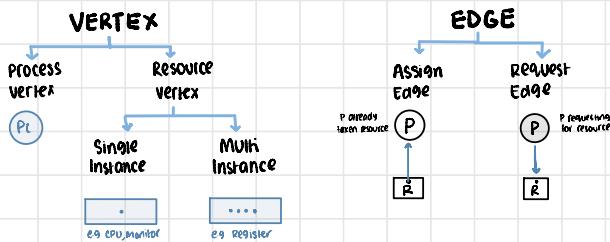
4. Circular Wait loop must exist



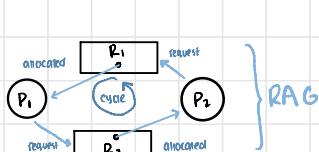
Resource Allocation Graph (RAG)

TO DETECT
DEADLOCKS

↳ whenever there is a deadlock or not, to represent we use RAG



SINGLE INSTANCE RAG



	Allocate		Request	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0

Availability = $(\frac{1}{2}, \frac{1}{2})$ → means no resource available

DEAD LOCK SITUATION

	Allocate		Request	
	R1	R2	R1	R2
P1	1	0	0	0
P2	0	1	0	0
P3	0	0	1	1

Availability = $(\frac{1}{2}, \frac{1}{2})$

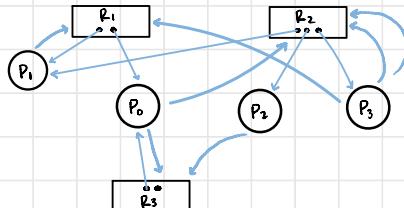
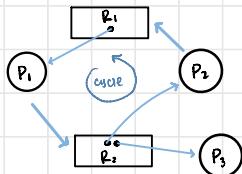
$(\frac{1}{2}, \frac{1}{2})$

$(\frac{1}{2}, \frac{1}{2})$

$P_1 \rightarrow P_2 \rightarrow P_3$

NO DEADLOCK

MULTI INSTANCE RAG



	Allocate		Request	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0
P3	0	1	0	0

Curr Availability = $(\frac{1}{2}, \frac{1}{2})$

$(\frac{1}{2}, \frac{1}{2})$

$(\frac{1}{2}, \frac{1}{2})$

$(\frac{1}{2}, \frac{1}{2})$

resources freed
 $(0, 1)$ = R1, R2 both 2 available
 since row 1 is free
 $(1, 1)$

$P_3 \rightarrow P_2 \rightarrow P_1$

	Allocated			Request		
	R1	R2	R3	R1	R2	R3
P0	1	0	1	0	1	1
P1	1	1	0	1	0	0
P2	0	1	0	0	0	1
P3	0	1	0	1	2	0

Curr Availability = $(0, 0, 1)$

$P_2 \rightarrow P_0 \rightarrow P_1 \rightarrow P_3$

(0, 1, 1)
(1, 1, 2)
(2, 2, 2)
(2, 3, 2)

Various methods to handle deadlock

1. Deadlock Ignorance (Ostrich method)

↳ just ignore deadlock as deadlock occurs ~~rarely~~

so why write an algo...

which will reduce performance → as we want more speed

if occurs just restart system

widely used



2. Deadlock Prevention

↳ before deadlock occurs try to find solution

by removing all 4 or any 1 of the necessary deadlock conditions

FAIL ANY ONE

1. Mutual Exclusion

↳ make all resources sharable

(can't be shared - for eg: monitor at most one thread can be in monitor)

3. Hold and Wait

↳ give a process all its demanding before doing anything

↳ if it can't get all resource then get none

→ Each process must wait for all its demands before doing anything else
→ If a process has all its demands then it can do anything else
→ If a process does not have all its demands then it can do nothing

2. No Preemption

↳ make preempted using a time quantum

→ The idea is to give each process a time quantum and then switch to another process after that time quantum has been completed.

4. Circular Wait

↳ order all resources

↳ assign each resource a priority

↳ make processes acquire resources in priority order

Resources
1. Printer
2. Scanner
3. CPU
4. Memory

4. Deadlock Detection & Recovery

↳ allow system to enter deadlock state

↳ detect deadlock

↳ recover deadlock

↳ kill the process

then check if still deadlock
if yes then repeat

↳ Resource Preemption

preempt a resource from a process

- 1. Scheduling a higher-priority process
- 2. Relocation - move only those shared processes for the deadlocked process
- 3. Reservation - seize process-share strategies picked by a user, handle conflicts or conflicts on your behalf

3. Deadlock Avoidance (Bankers Algorithm)

↳ when we give resources to a process, we check if it's safe or not

1. Process Initiation Denial

→ not optimal as worst assumes the worst

↳ process only starts if max claims

of all current + new processes can be met

2. Resource Allocation Denial

↳ aka Bankers Algorithm

1. Safe state → no deadlocks

2. Unsafe state → possibility of deadlock

3. Avoidance → ensures system will never enter unsafe state

Bankers Algo

used by
Deadlock Avoidance
Deadlock Detection

Process	Allocation			Max need	Current Available	Remaining Need			
	CPU	Memory	Printer			A	B	C	A
Resources	A	B	C	A	B	C	A	B	C
P ₁	0	1	0	7	5	3	3	3	2
P ₂	2	0	0	3	2	2	5	3	2
P ₃	3	0	2	9	0	2	7	4	3
P ₄	2	1	1	4	2	2	7	5	3
P ₅	0	0	2	5	3	3	10	5	5
	7	2	5				10	5	7

↑ 3 diff type resources
A=10, B=5, C=7 Freed allocated resources

↑ Total
→ allocated
↓ max need - allocation
Remaining Need

↓ fill this section last using remaining need

Safe sequence: P₂ → P₄ → P₁ → P₃ → P₅ → can have multiple sequences

Livelock

↳ Similar to deadlock except

processes never realise they are blocked
they change their regard

- For example, consider two processes each waiting for a resource the other has but waiting in a non-blocking manner
- When each learns that cannot continue they release their held resource and sleep for 30 seconds. Then they acquire their original resource followed by trying to the resource the other process held; they left, was occupied.
- Since both processes are trying to acquire (not hold), this is a livelock.

Memory Management

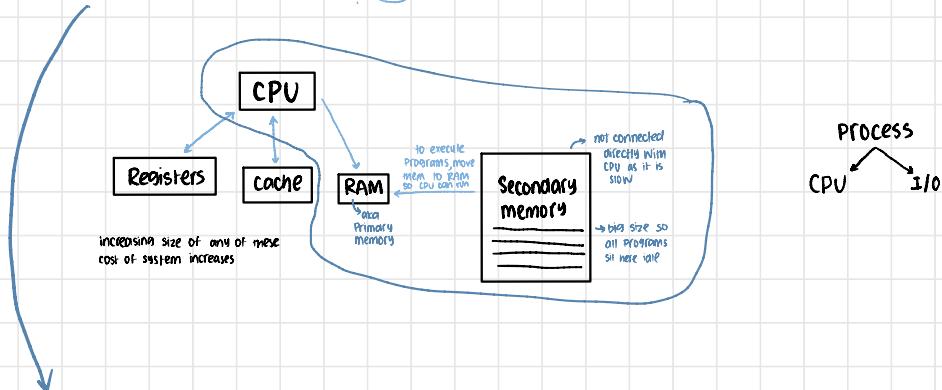
5.2

HOW DEGREE OF MULTIPROGRAMMING RELATES TO MEMORY MANAGEMENT

HOW TO MANAGE ALL RESOURCES IN EFFICIENT WAY BY

Method of managing Primary memory as majority of memory is RAM

why?



Degree of Multiprogramming

- moving more than 1 process in RAM → so whenever CPU needs process to execute, if we have plenty processes available so CPU not idle
- try to keep as many processes in RAM
- greater RAM, greater no. of processes in RAM
- greater degree, greater CPU utilization hence increase RAM

OS

ALLOCATION/DEALLOCATION

SECURITY

E.9

RAM 4mb

Process 4mb

$$\frac{4}{4} = 1 \text{ process}$$



time for I/O operation
 $k = 70\%$

$$\text{CPU Utilization} = (1 - k) \rightarrow 30\%$$

RAM 8mb

Process 4mb

$$\frac{8}{4} = 2 \text{ processes}$$



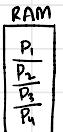
time for I/O operation
 $k = 70\%$

$$\text{CPU Utilization} = (1 - k^2) \rightarrow 76\% \quad ?$$

RAM 16mb

Process 4mb

$$\frac{16}{4} = 4 \text{ processes}$$



time for I/O operation
 $k = 70\%$

$$\text{CPU Utilization} = (1 - k^4) \rightarrow 94\% \quad ?$$

Address Binding

The binding of instructions and data to memory address

1. Compile time

- ↳ if you know at compile time where the process will reside in memory then **absolute code** can be generated
- ↳ if starting location changes, then necessary to recompile this code

2. Load time

- ↳ if you don't know at compile time where the process will reside in memory then **relocable code** must be generated
- ↳ if starting location changes, then only reload user code
- ↳ final binding delayed until **load time**

3. Execution time

- ↳ if process can be moved during execution from one memory segment to another then binding delayed until **run time**

Static linking – system libraries and program code combined by the loader into the binary program image

Dynamic linking – linking postponed until execution time

Small piece of code, **stub**, used to locate the appropriate memory-resident library routine

Stub replaces itself with the address of the routine, and executes the routine

Dynamic linking is particularly useful for libraries

Static Linking vs Dynamic Linking

Windows:
.lib vs .dll files

Linux:
.a vs .so files

Logical address (LA)

- ↳ aka virtual address
- ↳ always generated by CPU

Physical address (PA)

- ↳ aka absolute address
- ↳ address seen by m/m

Protection of Process Address space

1. make sure each process has a separate memory space
as protects processes from each other
2. protection is provided by these two registers
 - ↳ base register: holds smallest physical memory address
 - ↳ limit register: specifies size of range
3. Only OS can load base and limit registers
 - ↳ prevents user programs from changing register contents

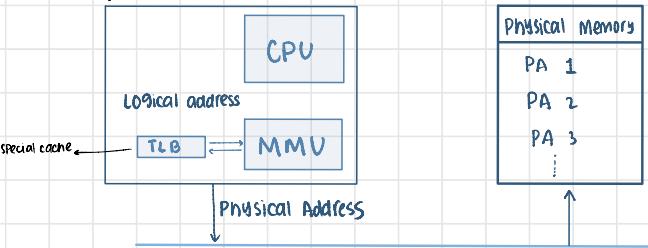
Relocation of Address

- *base register aka relocation register
- ↳ value of relocation register is added to every address generated by a user process

Memory management unit (MMU)

- ↳ converts LA into PA

- ↳ page table



5. ~ \rightarrow keep as many no. of processes in RAM \rightarrow Process in RAM are in Ready state

Memory Management Techniques

NO SPAN CONSECUTIVE

CONTINUOUS

static
Fixed Partition

dynamic
Variable Partition

SPAN NOT CONSECUTIVE

NON CONTINUOUS

Paging

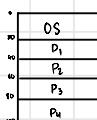
Multi-level
Paging

Invalid
Paging

Segmentation

Segmented
Paging

memory block
(RAM)

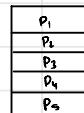


Process

Used some portion of in process
 \rightarrow in one partition and
the other in another

stored line wise

Process Partitioned



memory block
(RAM)



\rightarrow already occupied
by some other process

stored
in random
locations

CONTINUOUS MEMORY ALLOCATION

1. FIXED PARTITIONING (static partition) \rightarrow EASY TO IMPLEMENT

- ↳ No. of partitions are fixed
- ↳ size of partition may or may not vary
- ↳ stored line wise

LIMITATIONS

1. Internal Fragmentation \rightarrow Wastage of memory

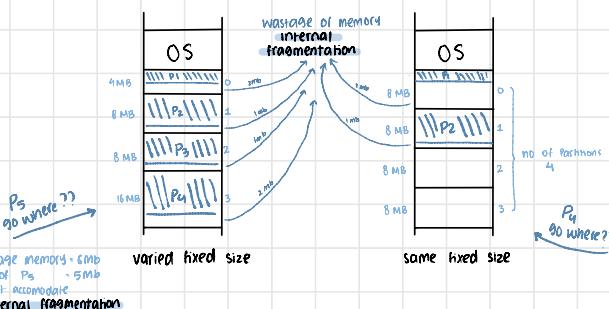
2. Limit in Process Size \rightarrow e.g. 16 MB as it's max size

3. Limitation on Degree of Multiprogramming

\rightarrow sum of all space = available space
but still apt. to accommodate
because of continuous mem.

can't accommodate more
than 4 processes as
no. of partitions are 4

4. External Fragmentation



Processes

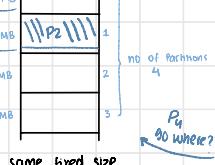
P₁ 2 Mb

P₂ 7 Mb

P₃ 7 Mb

P₄ 14 Mb

P₅ 5 Mb

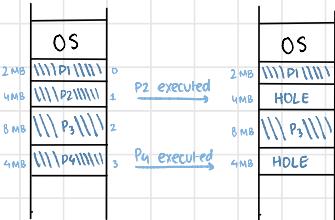


2 VARIABLE PARTITIONING (dynamic Partition)

- PROS
 - ↳ NO internal fragmentation
 - ↳ NO Limit on Degree of Multiprogramming
 - ↳ NO Limit in Process size
 - ↳ stored line wise

Processes

P ₁	2 Mb
P ₂	4 Mb
P ₃	8 Mb
P ₄	4 Mb



LIMITATIONS

1. External Fragmentation → sum of all space = available space but still not able to accommodate because of continuous memory
2. Allocation/Deallocation → complex as memory allocated dynamically at runtime holes

now P₆ = 8 Mb comes where to put it?

External Fragmentation

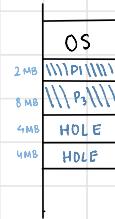
SOLUTION

COMPACTON →

- ↳ removes external fragmentation
- ↳ removes all holes together

CONS

- ↳ will have to stop all processes to move
- ↳ will have to move by copying which will take a lot of time



DYNAMIC

4 ALGOs to allocate processes in hole

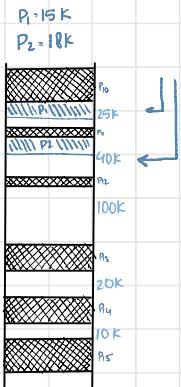
First-Fit most convenient

- ↳ allocate the first hole that is big enough
- ↳ always search from top

PRO

Fast

CON



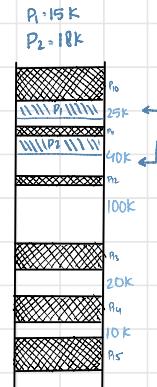
Next-Fit

- ↳ same as first fit but start search from last allocated hole

PROS

- ↳ faster than first fit as searches from last allocation

CONS



Best-Fit

- ↳ allocate the smallest that is big enough

PROS

- ↳ less internal fragmentation

CONS

- ↳ creates very tiny hole so former processes can't fit in it

slow → as searches entire list

Worst-Fit

allocate the largest that is big enough

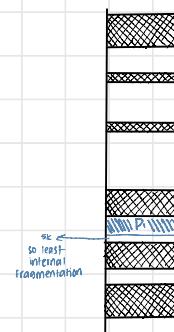
PROS

- ↳ creates big holes so former processes can fit in it

CONS

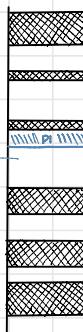
- ↳ great internal fragmentation
- slow → as searches entire list

P1 = 15 K



so less internal fragmentation

P1 = 15 K



so greatest internal fragmentation

keep some portion of a process
in one part and
the other in another

NON CONTINUOUS MEMORY ALLOCATION

↳ holes created dynamically

↳ dividing processes in run time

TIME CONSUMING

SOLUTION

process divided before hand
in secondary memory

main memory divided
before hand

PAGING

↳ removes external fragmentation

↳ divide processes into pages

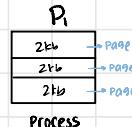
↳ each and every process has a PT

↳ bring some of those pages into M/M (which is divided into frames)

↳ PT will be in M/M

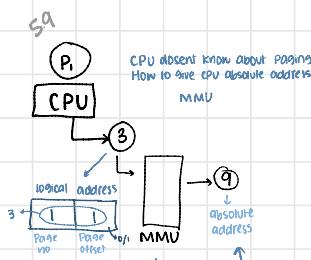
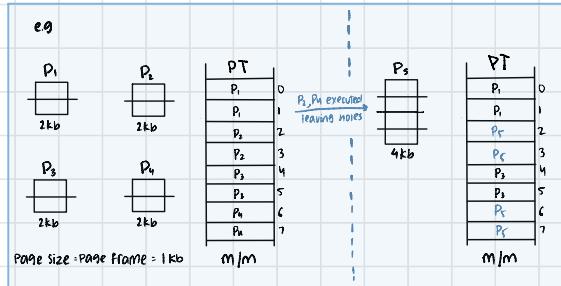
↳ when CPU demands a page/process

↳ search PT to find page frame



Page Size = frame size

as pages go inside frame
so perfect fit



Process size: 4b
Page size: 2b
no of pages: $\frac{4}{2} = 2$

P ₁	0 1 2 3	bytes
P ₁	0 1 2 3	bytes
M/M	0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	bytes
M/M	0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	bytes

M/M size: 16b

frame size = 2b

no of frames: $\frac{16}{2} = 8$ frames

5.10

represents size of process

$$\text{LAS} = 4\text{GB} = 2^3 \times 2^{10} \cdot 2^{32} \rightarrow \text{LA}$$

$$\text{PAS} = 64\text{MB} = 2^4 \times 2^{10} \cdot 2^{26} \rightarrow \text{PA}$$

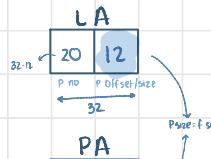
$$\text{Page size} = 4\text{KB} = 2^2 \times 2^{10} \cdot 2^9$$

$$\text{a) No of pages: } 32-12=20 = 2^{10}$$

$$\text{b) No of frames: } 26-12=14 = 2^4$$

$$\text{c) No of addresses in Page table: } 2^{20} \text{ as no of pages: } 2^{10}$$

$$\text{d) Size of Page table: } 2^{20} \times 14 =$$



Page size = frame size as pages fit inside frames

No of entries in PT = no of pages

PT size = no of entries \times size of entry $\xrightarrow{\text{no of frames bits}}$

Word = Byte (unless mentioned)

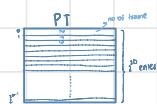
$$2^{32} = 32 \quad \begin{smallmatrix} 32 \\ \text{bits} \end{smallmatrix}$$

B: 8 bits \rightarrow so no need to convert

$$K = 2^{10 \text{ bits}}$$

$$M = 2^{10 \text{ bits}}$$

$$G = 2^{30 \text{ bits}}$$



5.11

$$\text{8) LA} = 7 \text{ bits}$$

$$\text{PA} = 6 \text{ bits}$$

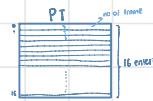
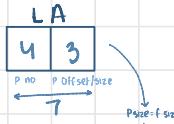
$$\text{Page size} = 8 \text{ words} = 2^3 = 3 \text{ bits}$$

$$\text{a) No of pages: } 7-3-4 = 2^4 = 16$$

$$\text{b) No of frames: } 6-3 = 3 = 2^3 = 8$$

$$\text{c) No of entries in PT: } 16$$

$$\text{d) Size of Page table: } 16 \times 3 =$$



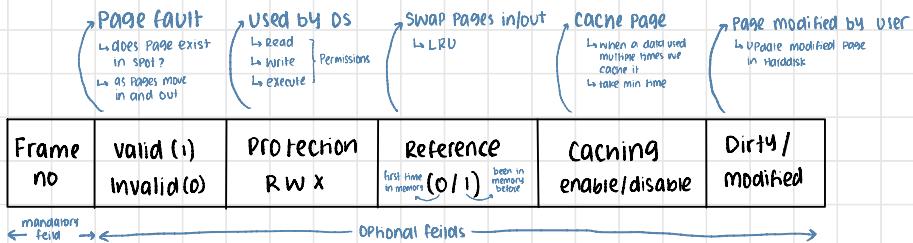
MEMORY MANAGEMENT UNIT (MMU)

5.12

↳ PAGE TABLES

↳ maps logical address to physical address

↳ every process has its own page table



5.13

2 level Paging / Hierarchical Paging

↳ if Page table size > page frame size

↳ so can't fit

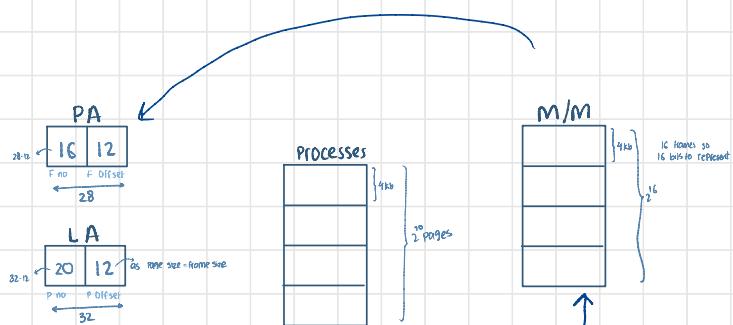
↳ so divide page table to smaller pages

$$PA/MM = 256 \text{ Mb} \rightarrow 28 \text{ bits}$$

$$LA = 4GB \rightarrow 32 \text{ bits}$$

$$\text{Frame size} = 4KB \rightarrow 2^{12} \text{ bits}$$

$$\text{Page Table Entry} = 2B$$



Page Table size: no of entries × size of entry

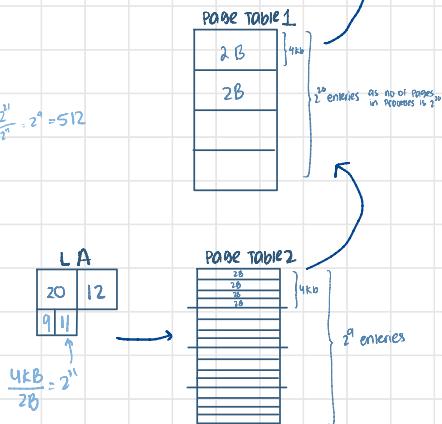
$$= 2^{20} \times 2B = 2MB$$

↳ can not fit in frame
so divide page table further into pages

Page Table size: no of entries × size of entry

$$= 2^9 \times 2B = 1KB$$

↳ can fit in frame



5.4 Inverted Paging

as m/m is limited, no of frame is limited

- Instead of each and every process has a PT,

all processes have only 1 PT \Rightarrow Global Page Table

no of entries in PT = no of frames

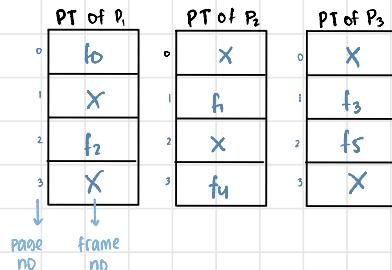
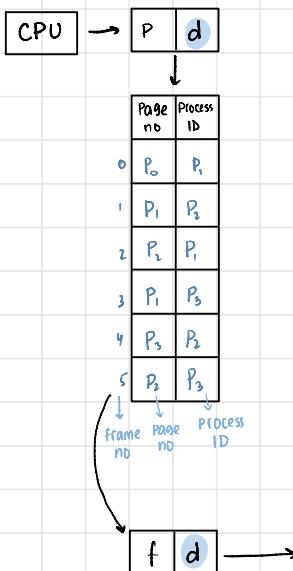
Normal Paging

- each and every process has a PT
- PT will be in m/m

no of entries in PT = no of pages

CON

memory is limited



5.5

$$8) \text{ VAS} = 32 \text{ bits}, \text{ PAGE SIZE} = 4 \text{ KB}, \text{ SYSTEM RAM} = 128 \text{ KB}$$

a) what is the ratio of VT and Inverted PT size if, table entry size = 4B

VA \rightarrow LA



no of VT entries: 2^{20}

VT size = $2^{20} \times 4$

IPt \rightarrow PA

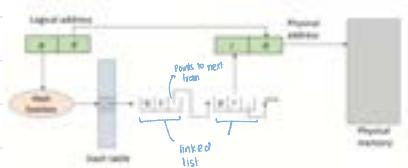


no of IPt entries = 2^5

IPt size = $2^5 \times 4$

$$\frac{2^{20} \times 4}{2^5 \times 4} = \frac{2^{16}}{1}$$

Hashed Page Table



Virtual Memory

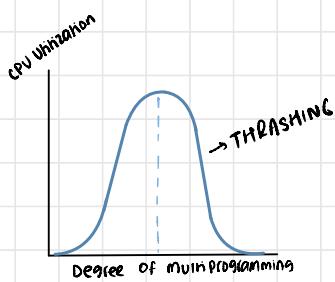
516

PAGE FAULTS

- ↳ bring some page of each process in RAM to increase DMR
- ↳ now some portion of all processes in RAM
- ↳ if CPU ask for a page that isn't in RAM → Page fault $\xrightarrow{\text{to fix}}$ Page fault service time
 - ↳ brings page from HDD to MM
 - ↳ take a lot of time
- ↳ so OS gets busy in servicing page fault
- ↳ Performance degrade

THRASHING

- ↳ page faults > page hit
- ↳ page fault service time → OS busy
- ↳ CPU remains idle for longer
- ↳ CPU utilization decreases



Degree of multi Programming

- ↳ no of processes in RAM increases
- ↳ CPU utilization increases

if CPU utilization too low

- ↳ increase degree of multiprogramming

TO AVOID THRASHING

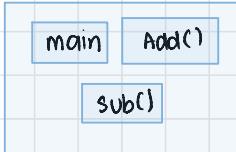
1. Increase size of M/M → difficult to achieve
 - ↳ don't take DMR to max level, decrease it a bit
2. Long term scheduler

Segmentation

- ↳ Process divided into parts/segments
- ↳ then put in M/m

PRO

- ↳ works from user point of view
- ↳ makes segments according to the code

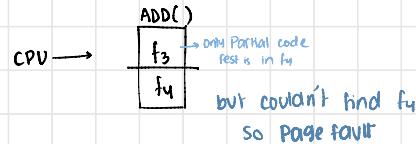


Paging

- ↳ Process divided into fixed sized pages
- ↳ then put in M/m

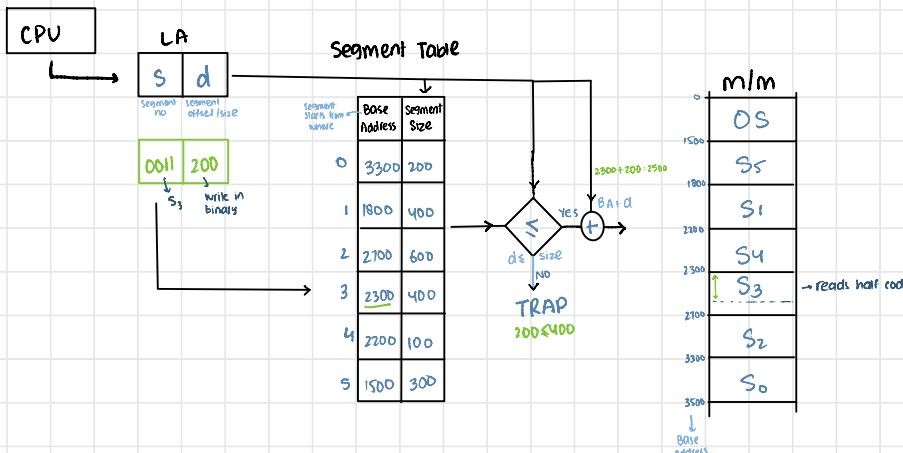
CON

- ↳ w/o knowing the what's the code we just divide in further pages



- ↳ CPU generates logical address
- ↳ will then asses what code it needs and tell where
- ↳ convert to physical address

LA → MMU → PA



5.10

VIRTUAL MEMORY

- ↳ gives illusion that a process whose size is larger than the size of m/m can be executed

↳ no limitation on

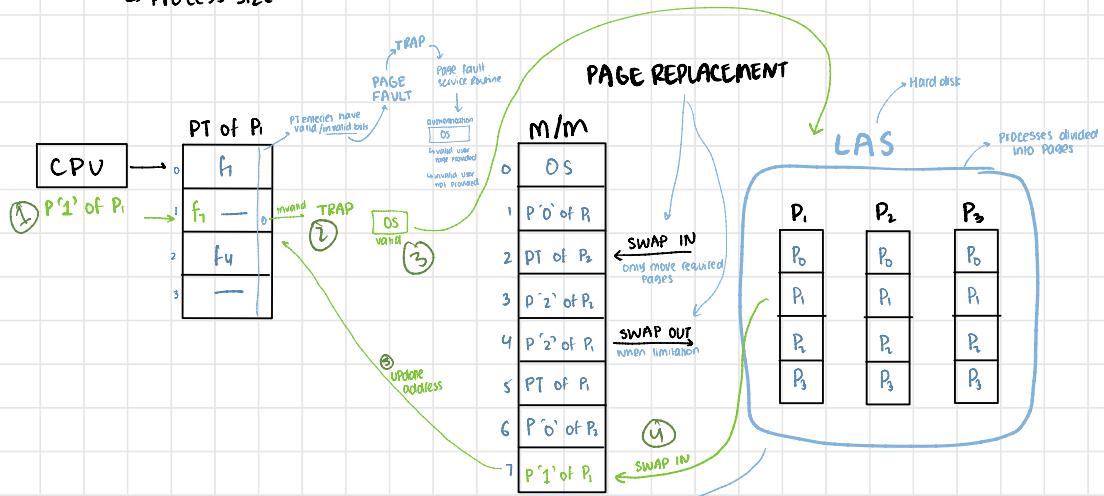
↳ no. of processes

↳ process size

↳ we currently use

LOCALITY OF REFERENCE

- ↳ when we bring a page in m/m we bring related pages along with it to



Demand Paging Steps

1. CPU generates
2. Page invalid \rightarrow Page fault
 - ↳ Trap generated
 - ↳ Page fault service routine
 - ↳ 3. OS authenticates
 - ↳ OS finds Page in LAS/Harddisk
 - 4. if frame free swap in the page to m/m
 - ↳ not free uses Page Replacement algorithm
 - 5. Updates address in PT
 - 6. CPU given control

User

OS

User

Demand Paging

- ↳ less I/O needed
- ↳ less memory needed
- ↳ faster response
- ↳ more users

(EMAT)

$$\text{effective memory access time} = (P)(\text{Page fault service time}) + (1-P)(\text{m/m access time})$$

Probability of page fault

mostly in milliseconds (ms)

no page fault

mostly in nano-seconds (nsec)

Page fault service routine

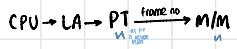
To compute the effective access time, we must know how much time is needed to service a page fault. A page fault causes the following sequence to occur:

1. Trap to the operating system.
2. Save the registers and process state.
3. Determine that the interrupt was a page fault.
4. Check that the page reference was legal, and determine the location of the page in secondary storage.
5. Create a read from the storage to a free frame.
 - a. Wait in a queue until the read request is serviced.
 - b. Wait for the device seek and/or latency time.
 - c. Begin the transfer of the page to a free frame.
6. While waiting, allocate the CPU core to some other process.
7. Receive an interrupt from the storage I/O subsystem (I/O completed).
8. Save the registers and process state for the other process (if step 6 is exercised).
9. Determine that the interrupt was from the secondary storage device.
10. Convert the page table and other tables to show that the desired page is now in memory.
11. Wait for the CPU core to be allocated to this process again.
12. Restore the registers, process state, and new page table, and then resume the interrupted instruction.

Page fault service routine with Page Replacement

1. Find the location of the desired page on secondary storage.
2. Find a free frame:
 - a. If there is a free frame, use it.
 - b. If there is no free frame, use a page-replacement algorithm to select a victim frame.
 - c. Write the victim frame to secondary storage (if necessary); change the page and frame tables.
3. Read the desired page into the newly freed frame; change the page and frame tables.
4. Continue the process from where the page-fault occurred.

↳ if PT is normal sized



$$\text{Total time} = 2n \rightarrow \text{ideal condition}$$

↳ if PT too big then
divided into multiple levels

Total time will increase

SOLUTION?

Translation Lookaside Buffer (TLB)

↳ to get faster memory use TLB / cache / buffer

as TLB is faster than RAM, so store PT in TLB

if no page fault

$$EMAT = HIT(TLB + m/m) + MISS(TLB + PT + m/m)$$

same as m/m

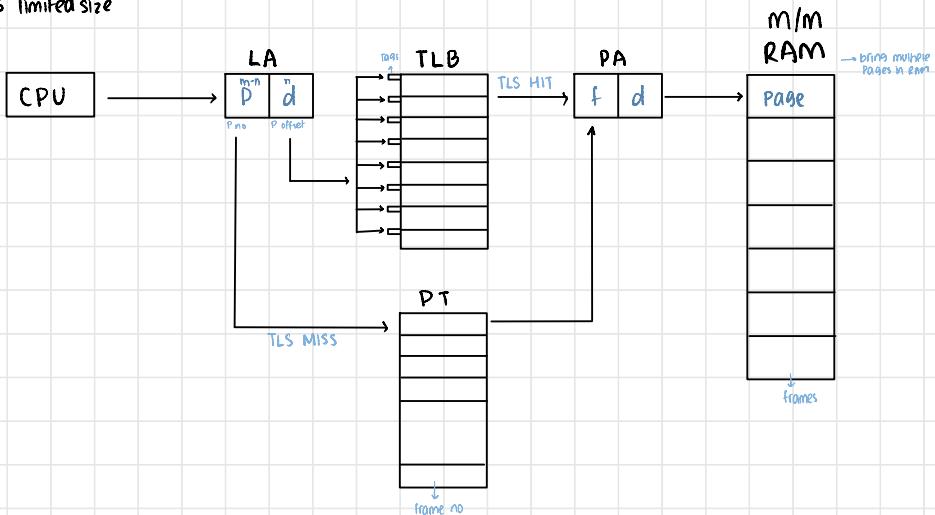
↳ used when Hits > Miss

↳ TLB HIT: found corresponding frame for page TLB $\xrightarrow{\text{frame no}} M/M$

↳ TLB Miss: then normal paging procedure TLB $\xrightarrow{\text{X}} PT \xrightarrow{\text{frame no}} M/M$

CON

↳ TLB has limited size



Q) Paging Scheme TLB

TLB Access time: 10 ns, M/M Access time: 50 ns

What is EMAT (in ns) if, HIT=90% and no page fault

$$EMAT = HIT(TLB + m/m) + MISS(TLB + PT + m/m)$$

$$= 0.9(10 + 50) + 0.1(10 + 50 + 50) = 65 \text{ ns}$$

5.22 PAGE REPLACEMENT

↳ In Virtual memory concept

↳ When M/M filled and need to bring new page in M/M

↳ Swap in, swap out

$$\text{Hit ratio} = \frac{\text{no. of hits}}{\text{no. of references}} \times 100$$

$$\text{Miss ratio} = \frac{\text{no. of miss}}{\text{no. of references}} \times 100$$

1. FIRST IN FIRST OUT (FIFO)

↳ When page fault occurs, replace page

which came in first

Q) Reference string = 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 1, 2, 0

f_3		1	1	1	X	0	0	0	3	3	3	3	2	2		
f_2		0	0	0	0	3	3	3	2	2	2	2	X	1	1	1
f_1	7	7	X	2	2	2	2	2	4	4	4	4	0	0	0	0
	*	*	*	X	HIT	*	*	*	*	*	*	*	HIT	*	*	HIT

↓
 Miss
 DS no space
 Replace → as if
 come in first!

$$\text{Page Hit} = 3 \quad \text{Hit ratio} = \frac{3}{15} \times 100 =$$

$$\text{Page fault} = 12 \quad \text{Miss ratio} = \frac{12}{15} \times 100 =$$

Q) Ref → 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

USING 3 FRAMES

f_3		3	3	3	2	2	2	2	X	4	4				
f_2		2	2	X	1	1	1	1	X	3	3	3			
f_1	1	1	X	4	4	4	5	5	5	5	5	5			
	*	*	*	*	*	*	*	*	HIT	HIT	*	*	HIT		

HITS = 3
MISS = 9

USING 4 FRAMES

f_4		4	4	4	4	4	4	4	4	3	3	3		
f_3		3	3	3	3	3	3	3	2	2	2	2		
f_2		2	2	2	2	2	X	1	1	1	X	5		
f_1	1	1	1	1	1	X	5	5	5	5	8	4	4	
	*	*	*	*	HIT	HIT	*	*	*	*	*	*	*	

$$\text{HITS} = 2 \rightarrow \text{Belady's Anomaly}$$

↳ increasing no of frames
increased no of page faults

2. OPTIMAL PAGE REPLACEMENT

↳ When page fault occurs, replace page

whose demand is very late in future

Q) Ref → 1, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 1, 0, 1



fu		2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	7	7	7
fs		1	1	1	1	x	4	4	4	4	4	x	1	1	1	1	1	1	1
ft		0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
h	7	7	7	7	7	3	3	3	3	3	3	3	3	3	3	3	3	3	3
*	*	*	*	*	HIT	*	HIT	*	HIT	HIT	HIT	HIT	*	HIT	HIT	HIT	*	HIT	HIT

HIT = 12

MISS = 8

3. LEAST RECENTLY USED (LRU)

→ COUNTING BASED PAGE REPLACEMENT ALGO

↳ When page fault occurs, replace page

which is least recently used in the past

Q) Ref → 1, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 1, 0, 1



fu		2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
fs		1	1	1	1	x	4	4	4	4	4	x	1	1	1	1	1	1	1
ft		0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
h	7	7	7	7	7	3	3	3	3	3	3	3	3	3	3	3	3	7	7
*	*	*	*	*	HIT	*	HIT	*	HIT	HIT	HIT	HIT	*	HIT	HIT	*	HIT	HIT	

HIT = 12

MISS = 8

5.1b

→ COUNTING BASED
PAGE REPLACEMENT ALGO

4. MOST RECENTLY USED (MRU)

↳ When page fault occurs, replace page

which is most recently used in the past

Q) Ref → 1, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 1, 0, 1
 ↑ recent

fu		2	2	2	2	2	2	3	0	3	2	2	x	0	0	0	0
fs		1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
ft		0	0	0	0	3	0	4	4	4	4	4	4	4	4	4	4
h	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1

* * * * HIT * * * HIT * * * HIT HIT HIT * HIT HIT HIT HIT

HIT = 8
MISS = 12

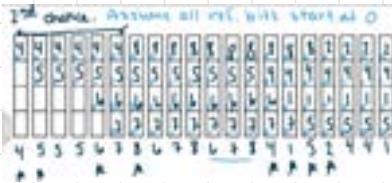
5. Second chance Algorithm

↳ uses reference bit

↳ FIFO

Ref	1	2	3	4	2	1	5	6	2	1	2	3	7	6	3	2	1	2	3	6
		30	30	30	10	10	10	20	20	20	20	20	20	20	30	30	30	30	30	
		20	20	20	20	20	20	60	60	60	60	60	60	60	60	60	60	60	60	
	10	10	10	40	40	40	50	50	50	50	50	50	50	50	20	20	20	20	20	

H H H



11.4.3 Enhanced Second-Chance Algorithm

We can enhance the second-chance algorithm by considering the reference bit and the modify bit (shown in Section 10.4.1) as an ordered pair. With these two bits, we have the following four possible states:

- 1. 00 neither recently used nor modified → swap page to replace
- 2. 01 recently used but not modified → swap page to avoid increase the page fault rate in memory page replacement
- 3. 10 recently used and modified → probably will be used again soon
- 4. 11 recently used and modified → probably will be used again soon, the page will be moved to the written buffer to minimize memory thrashing from its replacement

Enhance Second Chance

The major difference between ESC algorithm and the simpler SC algorithm is that here we give preference to those pages that have been modified in order to reduce the number of I/Os required.

COPY ON WRITE (COW)

↳ instead share pages

↳ Pages that can be modified COW marked

↳ if a P/C wants to write on that shared page

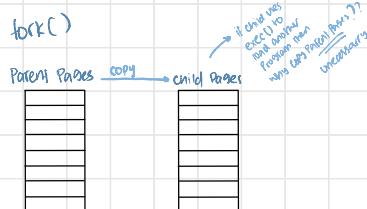
↳ a copy of that page is made as it is modified

access to old page lost

↳ P/C share pages that are unmodified

PRO

↳ efficient Process creation → all pages aren't duplicated
only the modified ones are



MODIFY/DIRTY bit

↳ modify bit with each frame

↳ modify bit = 1

↳ If changes are done

↳ update the changes in HD/
write back in HD

	v	j	modify
0	5	1	1 → if no
1			
2			
3			
4			
5			
6			
7			
8			
9			
10			
11			
12			
13			
14			
15			

↳ modify bit = 0

↳ no changes are done

↳ no need to update as copy

already in HD

Page and Frame Replacement Algorithms

Frame-allocation algorithm

- How many frames to give each process

- Which frames to replace

Page-replacement algorithm

- Min lowest page-fault rate on both first access and re-access

PROS

↳ reduces Page fault-time

↳ reduce page transfers overhead → at only modified pages are written back

Allocation of Frames

- Each process needs **minimum** number of frames
- **Maximum** of course is total frames in the system
- Two major allocation schemes
 - fixed allocation
 - priority allocation

Fixed Allocation

Equal allocation - For example, if there are 100 frames (after allocating frames for the OS) and 5 processes.

- Keep some as free frame buffer pool

- Equal Allocation = Number of frames / number of processes
- EA = 100 / 5 = 20 frames / process

Proportional Allocation

- Proportional allocation - Allocate according to the ~~size of~~ of processes

Dynamic as degree of multiprogramming, process sizes change

$$\begin{aligned} n &= \text{no. of processes} \\ B &= \Sigma n_i \\ m &= \text{total no. of frames} \\ a_i &= \text{allocation for } p_i = \frac{n_i}{m} \times m \end{aligned}$$
$$\begin{aligned} n_1 &= 64 \\ n_2 &= 39 \\ n_3 &= 127 \\ n_4 &= \frac{10}{137} \times 137 = 10 \\ n_5 &= \frac{127}{137} \times 137 = 127 \end{aligned}$$

Priority Allocation

- Use a proportional allocation scheme using **priorities** rather than size
- If process P_j generates a page fault,
 - select for replacement one of its frames
 - select for replacement a frame from a process with lower priority number

Global vs. Local Allocation

- **Global replacement** - process selects a replacement frame from the set of all frames; one process can take a frame from another

- But then process execution time can vary greatly
- But greater throughput so more common



- **Local replacement** - each process selects from only its own set of allocated frames

- More consistent per-process performance
- But possibly underutilized memory



Non-Uniform Memory Access (NUMA)

(Contd.) performance comes from allocating memory "close to" the CPU on which the thread is scheduled

- And modifying the scheduler to schedule the thread on the same system board when possible
- Solved by Solaris by creating **lgroups**
 - Structures to track CPU / Memory low latency groups
 - Used my scheduler and pages
- When possible schedules all threads of a process and allocates memory for that process within the lgroup

Page-Buffering Algorithms

- Keep a pool of free frames, always
 - Then frame available when needed, not found at fault time
 - Read page into free frame and select victim to evict and add to free pool
 - When convenient, evict victim
- Possibly, keep list of modified pages
 - When backing store otherwise idle, write pages there and set to non-dirty
- Possibly, keep free frame contents intact and note what is in them
 - If reference again before reused, no need to load contents again from disk

Applications and Page Replacement

- All of these algorithms have OS guessing about future page access
- Some applications have better knowledge - i.e. databases
- Memory intensive applications can cause double buffering
- Operating system can give direct access to the disk, getting out of the way of the applications

turn around time: completion - arrival time

Waiting time: Total waiting time - ms process executed - Arrival time

Preemptive

Waiting time = turn around time - burst time

↓
non preemptive

AMDAHL'S LAW

↳ speed up in latency of a task execution

$$\text{Speed Up} \leq \frac{1}{S + \frac{\sum_{i=1}^N t_i}{N}}$$

serial \rightarrow parallel
 N \rightarrow no. of cores

Page Size = frame size as pages go inside frames

no of entries in PT = no of pages

PT size = no of entries \times size of entry \rightarrow no of frames bits

word = byte (unless mentioned)

$$2^{32 \text{ bits}} = 32 \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline \end{array}$$

B = 8 bits \rightarrow so no need to convert

$$K = 2^{10 \text{ bits}}$$

$$M = 2^{20 \text{ bits}}$$

$$G = 2^{30 \text{ bits}}$$

if no page fault

$$\text{EMAT} = \text{HIT}(\text{TLB} + m/m) + \text{MISS}(\text{TLB} + \text{PT} + m/m)$$

same as m/m

(EMAT)

$$\text{effective memory access time} = (P)(\text{page fault service time}) + (1-P)(m/m \text{ access time})$$

Probability of Page Fault

mostly in milliseconds

no page fault

mostly in nano seconds

NO
TALKING

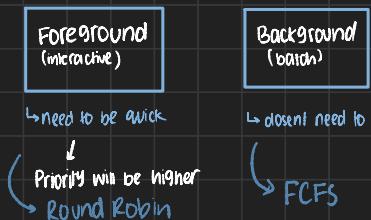
CLASS SAY

BAMIR

Multilevel Queue Scheduling

scheduling algorithms for situations in which processes are classified in different groups

E.g.



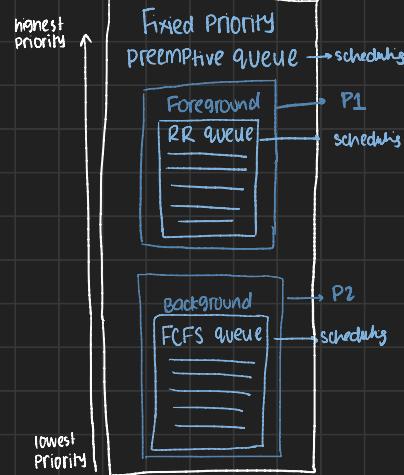
↳ need to be quick

↳ Priority will be higher
Round Robin

↳ doesn't need to be as fast

↳ FCFS

- ↳ partitions ready queue into several separate queues
- ↳ scheduling takes place within and among these queues
- ↳ each process permanently assigned to 1 queue
based on size, priority, type
- ↳ each queue has its own scheduling algorithm (FCFS, SJF, RR...)



Multilevel Feedback Queue Scheduling

- ↳ allows processes to move between queues
- ↳ separate processes according to their CPU bursts
- ↳ if a process uses too much CPU time
then moved to lower priority queue (so everybody gets a chance)
- ↳ I/O and interactive processes are in higher priority queues
as they need quick responses
- ↳ a process in lower priority queue for too long
may be moved to higher priority queue (prevents starvation)

Parameters

- ↳ no. of queues
- ↳ scheduling algo for each queue
- ↳ method to upgrade process to higher priority
- ↳ method to demote process to higher priority
- ↳ method to determine queue when process demoted/upgraded

PROCESS

- ↳ A program in execution
- ↳ can have 1 or more threads



Threads

- ↳ unit of execution Within a Process

- ↳ It compromises of

1. Thread ID
2. Program Counter
3. Register set
4. Stack

- ↳ shares code section, data section, os resources with other threads of same processes

	Process	Thread
1. No.	Process is heavier weight on resource utilization.	Thread is lighter weight on resource utilization than a process.
2.	Process controlling access interaction with operating system.	Thread switching does not need to interact with operating system.
3.	In multiprocess environment, each process executes the same code but has its own memory and file resources.	All threads can share same set of dependent files, code, processes.
4.	If one process in isolated then no other process can execute until that thread (process) is terminated.	One core thread is isolated from switching, switched thread in other core task can run.
5.	Multiprocess environment using threads can reuse resources.	Advantage: Intra-thread processing over inter-thread communication.
6.	No multiprocess environment each process represents the responsibility of that process.	One thread can read, write or change another thread's state.

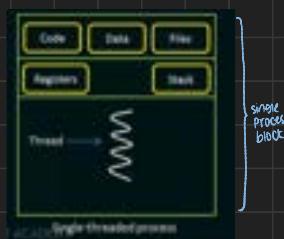
MULTITHREAD

- ↳ performs multiple tasks at a time



Single Thread

- ↳ performs one task at a time



MULTI-THREADING BENEFITS

- ↳ **Responsiveness:** may continue execution if process is blocked / performing lengthy operation

dedicated threads for handling user events

- ↳ **Resource Sharing:** threads share memory and resources of a process → allows diff threads activity within same address space
allocating memory/resources in cores ↗ easier than shared memory/message passing

thread switching has lower overhead than context switching

- ↳ **Economy:** cheaper than process creation as threads share resources

- ↳ **Scalability:** utilization of multiple cores for parallel execution

increases concurrency